

Exceptional Control Flow: Exceptions and Processes

CSE4100: Multicore Programming

Sungyong Park (PhD)

Data-Intensive Computing and Systems Laboratory (DISCOS)

<https://discos.sogang.ac.kr>


Office: R908A, E-mail: parksy@sogang.ac.kr

Control Flow

■ Processors do only one thing:

- From startup to shutdown, a CPU simply reads and executes (interprets) a sequence of instructions, one at a time
- This sequence is the CPU's *control flow* (or *flow of control*)

Physical control flow



Altering the Control Flow

- **Up to now: two mechanisms for changing control flow:**

- Jumps and branches
- Call and return

React to changes in *program state*

- **Insufficient for a useful system:**

Difficult to react to changes in *system state*

- Data arrives from a disk or a network adapter
- Instruction divides by zero
- User hits Ctrl-C at the keyboard
- System timer expires

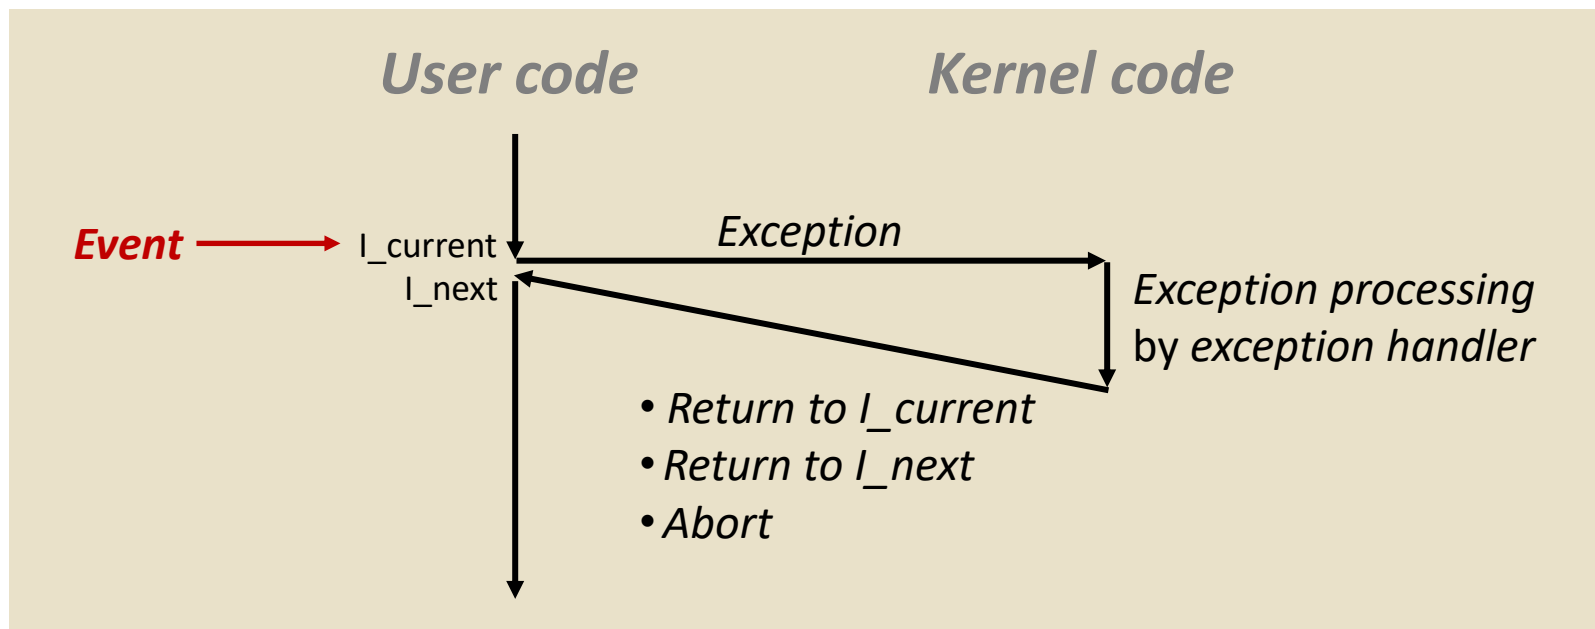
- **System needs mechanisms for “exceptional control flow”**

Exceptional Control Flow

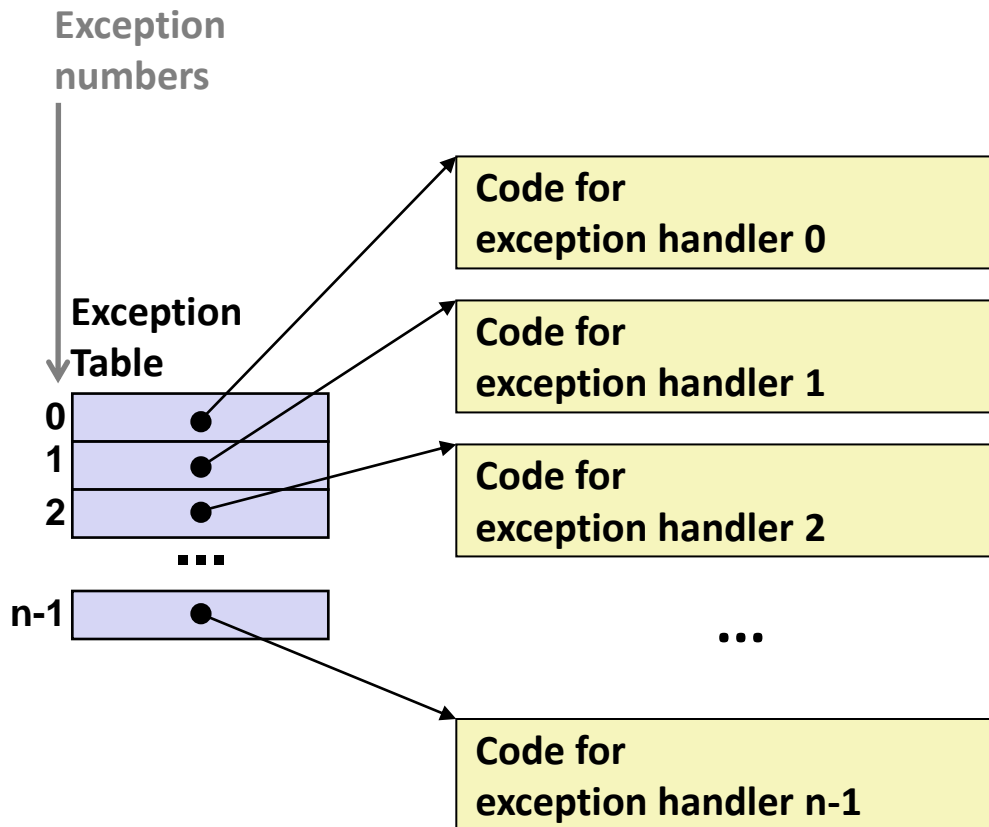
- **Exists at all levels of a computer system**
- **Low level mechanisms**
 - 1. **Exceptions**
 - Change in control flow in response to a system event (i.e., change in system state)
 - Implemented using combination of hardware and OS software
- **Higher level mechanisms**
 - 2. **Process context switch**
 - Implemented by OS software and hardware timer
 - 3. **Signals**
 - Implemented by OS software
 - 4. **Nonlocal jumps**: `setjmp()` and `longjmp()`
 - Implemented by C runtime library

Exceptions

- An **exception** is a transfer of control to the OS *kernel* in response to some *event* (i.e., change in processor state)
 - Kernel is the memory-resident part of the OS
 - Examples of events: Divide by 0, arithmetic overflow, page fault, I/O request completes, typing Ctrl-C



Exception Tables



- Each type of event has a unique exception number k
- k = index into exception table (a.k.a. interrupt vector)
- Handler k is called each time exception k occurs

Classes of Exceptions

■ Interrupts, Traps, Faults, and Aborts

Class	Cause	Async/sync	Return behavior
Interrupt	Signal from I/O device	Async	Always returns to next instruction
Trap	Intentional exception	Sync	Always returns to next instruction
Fault	Potentially recoverable error	Sync	Might return to current instruction
Abort	Nonrecoverable error	Sync	Never returns

Figure 8.4 **Classes of exceptions.** Asynchronous exceptions occur as a result of events in I/O devices that are external to the processor. Synchronous exceptions occur as a direct result of executing an instruction.

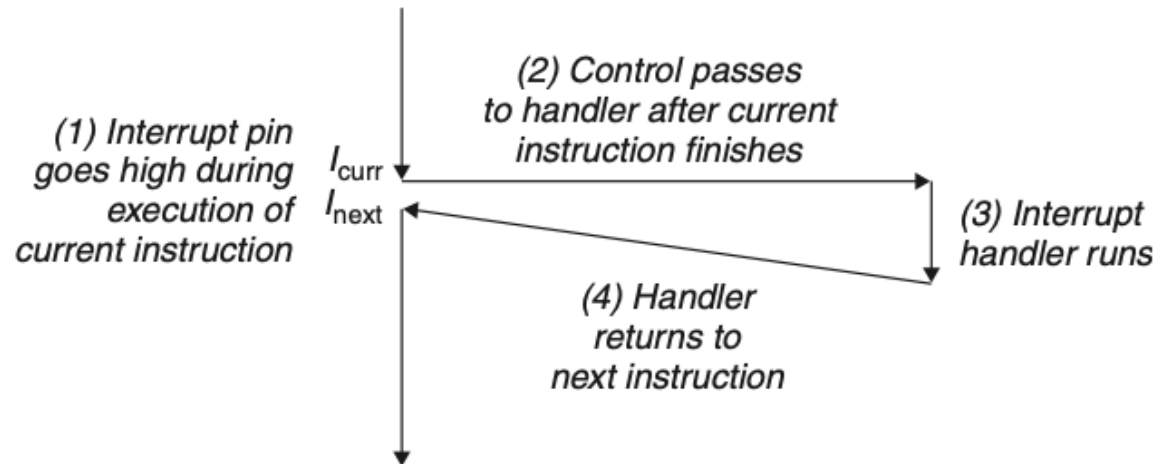
Asynchronous Exceptions (Interrupts)

- **Caused by events external to the processor**
 - Indicated by setting the processor's *interrupt pin*
 - Handler returns to “next” instruction

Figure 8.5

Interrupt handling.

The interrupt handler returns control to the next instruction in the application program's control flow.



Asynchronous Exceptions (Interrupts)

■ Examples:

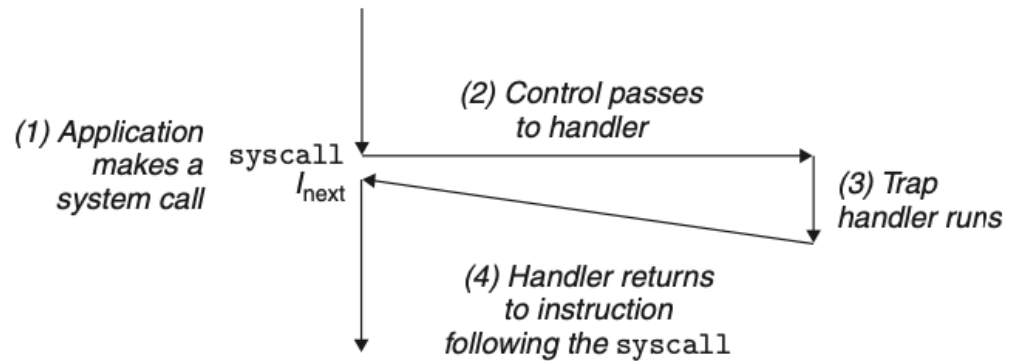
- Timer interrupt
 - Every few ms, an external timer chip triggers an interrupt
 - Used by the kernel to take back control from user programs
- I/O interrupt from external device
 - Arrival of a packet from a network
 - Arrival of data from a disk

Synchronous Exceptions

- **Caused by events that occur as a result of executing an instruction:**
 - ***Traps***
 - Intentional
 - Examples: ***system calls***, breakpoint traps, special instructions
 - Returns control to “next” instruction
 - ***Faults***
 - Unintentional but possibly recoverable
 - Examples: page faults (recoverable), protection faults (unrecoverable), floating point exceptions
 - Either re-executes faulting (“current”) instruction or aborts
 - ***Aborts***
 - Unintentional and unrecoverable
 - Examples: illegal instruction, parity error, machine check
 - Aborts current program

Figure 8.6

Trap handling. The trap handler returns control to the next instruction in the application program's control flow.



Synchronous Exceptions (Traps)

Examples: System calls

System Call Example: Opening File

- User calls: `open(filename, options)`
- Calls `__open` function, which invokes system call instruction `syscall`

0000000000e5d70 <__open>:

...

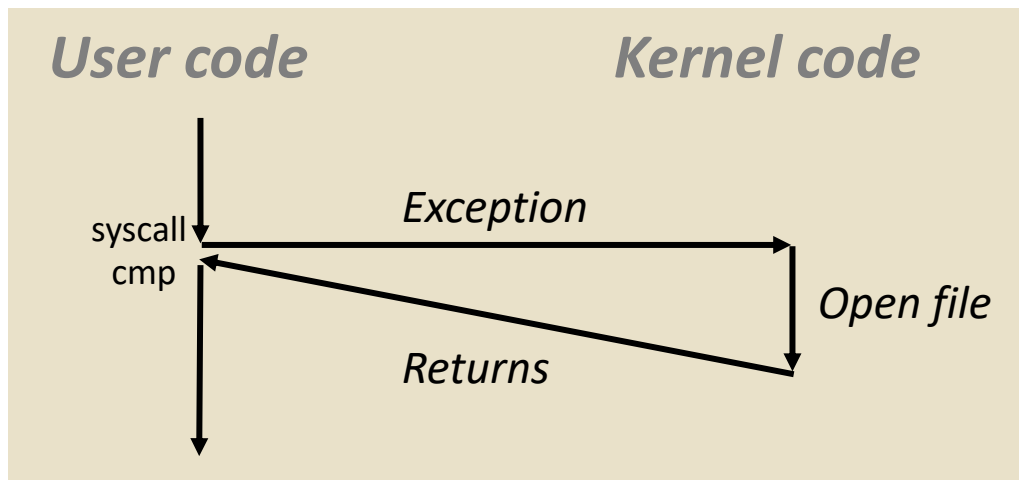
e5d79: b8 02 00 00 00 mov \$0x2,%rax # open is syscall #2

e5d7e: 0f 05 syscall # Return value in %rax

e5d80: 48 3d 01 f0 ff ff cmp \$0xfffffffffff001,%rax

...

e5dfa: c3 retq



- `%rax` contains syscall number
- Other arguments in `%rdi`, `%rsi`, `%rdx`, `%r10`, `%r8`, `%r9`
- Return value in `%rax`
- Negative value is an error corresponding to negative `errno`

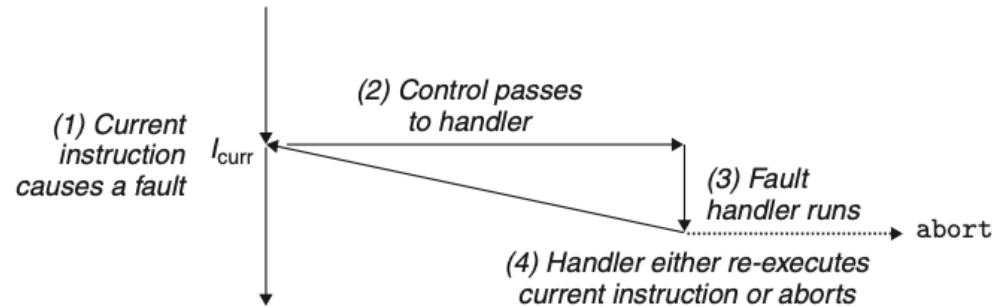
System Call Number

- Each x86-64 system call has a unique ID number
- Examples:

<i>Number</i>	<i>Name</i>	<i>Description</i>
0	read	Read file
1	write	Write file
2	open	Open file
3	close	Close file
4	stat	Get info about file
57	fork	Create process
59	execve	Execute a program
60	_exit	Terminate process
62	kill	Send signal to process

Figure 8.7**Fault handling.**

Depending on whether the fault can be repaired or not, the fault handler either re-executes the faulting instruction or aborts.



Synchronous Exceptions (Fault Handling)

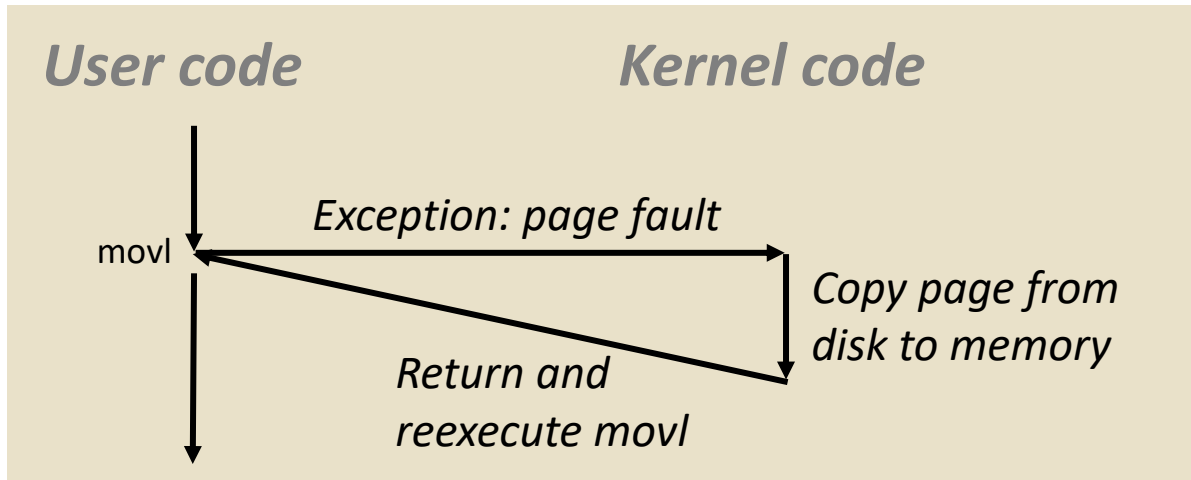
Examples: Page faults

Fault Example: Page Fault

- User writes to memory location
- That portion (page) of user's memory is currently on disk

```
int a[1000];  
main ()  
{  
    a[500] = 13;  
}
```


```
80483b7:      c7 05 10 9d 04 08 0d  movl    $0xd,0x8049d10
```



Fault Example: Invalid Memory Reference

```
int a[1000];  
main ()  
{  
    a[5000] = 13;  
}
```


80483b7: c7 05 60 e3 04 08 0d movl \$0xd,0x804e360



- Sends **SIGSEGV** signal to user process
- User process exits with “segmentation fault”

Figure 8.8

Abort handling. The abort handler passes control to a kernel **abort** routine that terminates the application program.




Synchronous Exceptions (Abort Handling)

Examples: Hardware errors such as parity errors that occur when DRAM or SRAM bits are corrupted.

Processes

- **Definition:** A *process* is an instance of a running program.
 - One of the most profound ideas in computer science
 - Not the same as “program” or “processor”
- **Process provides each program with two key abstractions:**
 - *Logical control flow*
 - Each program seems to have exclusive use of the CPU
 - Provided by kernel mechanism called *context switching*
 - *Private address space*
 - Each program seems to have exclusive use of main memory.
 - Provided by kernel mechanism called *virtual memory*




Control Flow

■ Processors do only one thing:


- From startup to shutdown, each CPU core simply reads and executes a sequence of machine instructions, one at a time *
- This sequence is the CPU's *control flow* (or *flow of control*)

Physical control flow



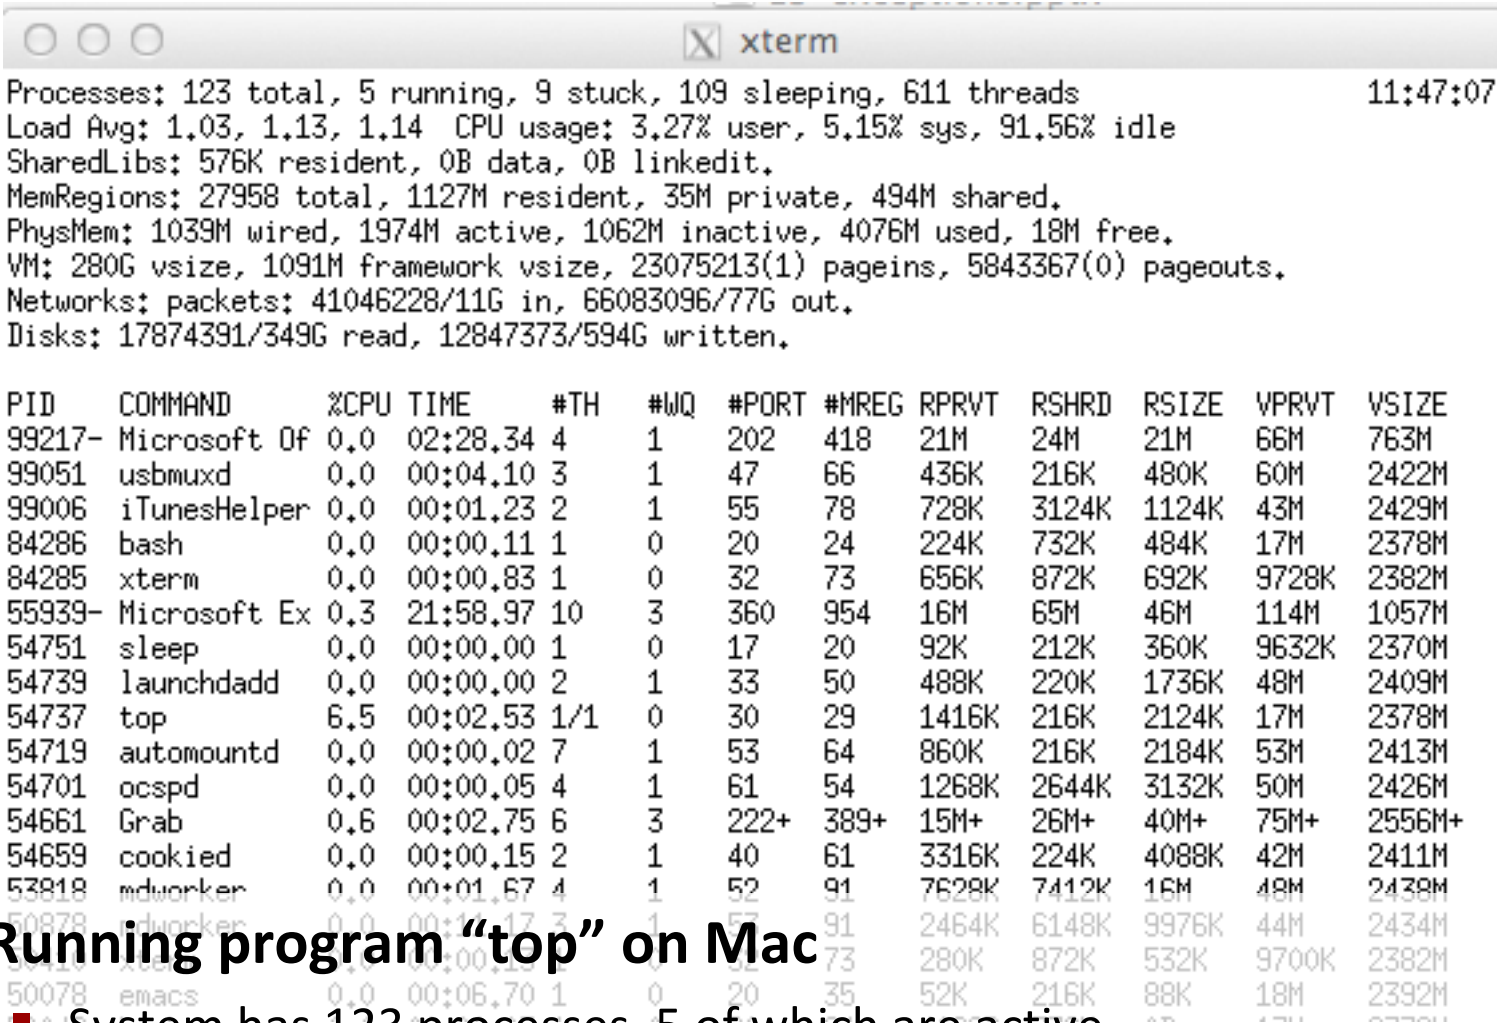
- * many modern CPUs execute several instructions at once and/or out of program order, but this is invisible to the programmer

Multiprocessing: The Illusion



- **Computer runs many processes simultaneously**
 - Applications for one or more users
 - Web browsers, email clients, editors, ...
 - Background tasks
 - Monitoring network & I/O devices

Multiprocessing Example



```

Processes: 123 total, 5 running, 9 stuck, 109 sleeping, 611 threads
Load Avg: 1.03, 1.13, 1.14  CPU usage: 3.27% user, 5.15% sys, 91.56% idle
SharedLibs: 576K resident, 0B data, 0B linkedit.
MemRegions: 27958 total, 1127M resident, 35M private, 494M shared.
PhysMem: 1039M wired, 1974M active, 1062M inactive, 4076M used, 18M free.
VM: 280G vsize, 1091M framework vsize, 23075213(1) pageins, 5843367(0) pageouts.
Networks: packets: 41046228/11G in, 66083096/77G out.
Disks: 17874391/349G read, 12847373/594G written.

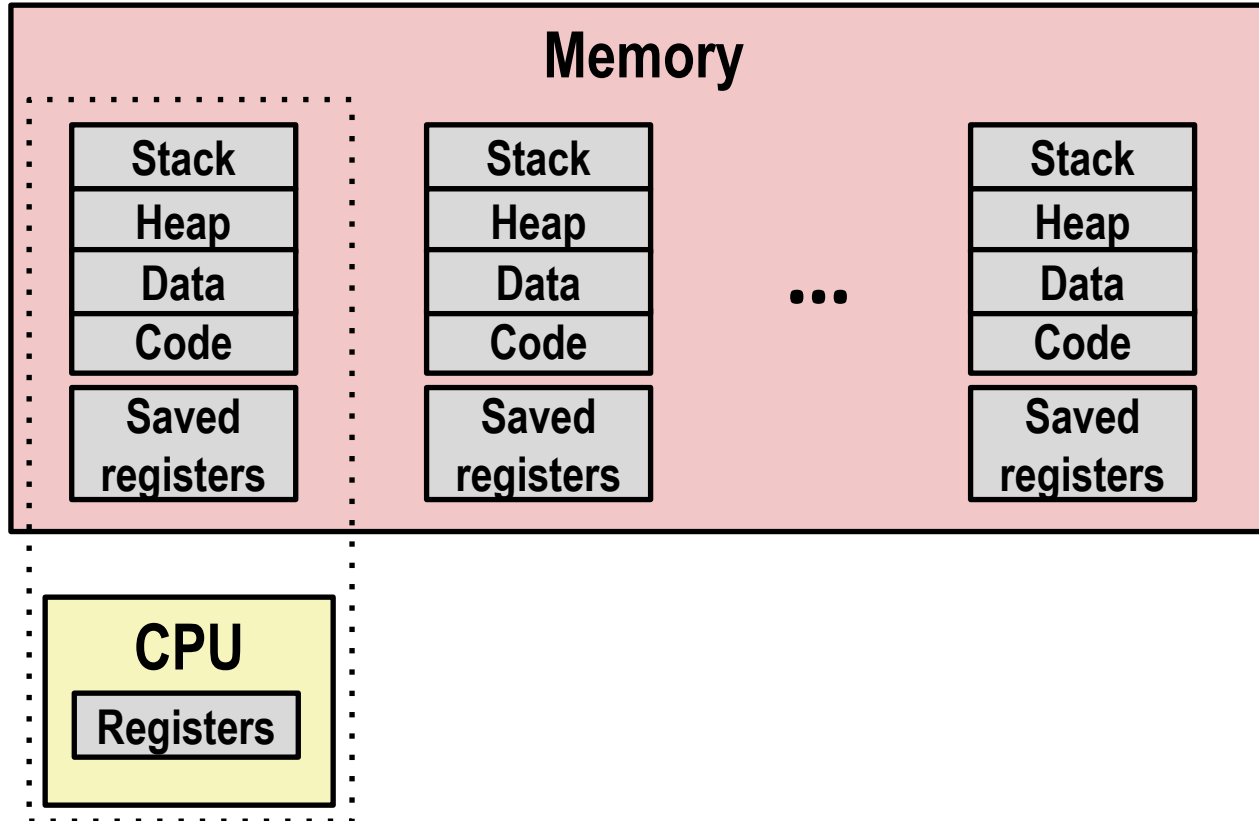
PID    COMMAND      %CPU TIME    #TH  #WQ  #PORT #MREG RPRVT  RSHRD  RSIZE  VPRVT  VSIZE
99217-  Microsoft Of 0.0  02:28.34  4    1    202  418   21M   24M   21M   66M   763M
99051   usbmuxd      0.0  00:04.10  3    1    47   66   436K  216K  480K  60M   2422M
99006   iTunesHelper 0.0  00:01.23  2    1    55   78   728K  3124K 1124K 43M   2429M
84286   bash         0.0  00:00.11  1    0    20   24   224K  732K  484K  17M   2378M
84285   xterm        0.0  00:00.83  1    0    32   73   656K  872K  692K  9728K 2382M
55939-  Microsoft Ex 0.3  21:58.97 10    3    360  954   16M   65M   46M   114M  1057M
54751   sleep        0.0  00:00.00  1    0    17   20   92K   212K  360K  9632K 2370M
54739   launchdadd  0.0  00:00.00  2    1    33   50   488K  220K  1736K 48M   2409M
54737   top          6.5  00:02.53 1/1    0    30   29  1416K  216K  2124K 17M   2378M
54719   automountd  0.0  00:00.02  7    1    53   64   860K  216K  2184K 53M   2413M
54701   ocsdpd       0.0  00:00.05  4    1    61   54  1268K  2644K 3132K 50M   2426M
54661   Grab         0.6  00:02.75  6    3   222+ 389+  15M+  26M+  40M+  75M+ 2556M+
54659   cookied      0.0  00:00.15  2    1    40   61  3316K  224K  4088K 42M   2411M
53818   mdworker     0.0  00:01.67  4    1    52   91  7628K  7412K  16M   48M   2438M
50978   mdworker     0.0  00:01.17  3    1    57   91  2464K  6148K 9976K 44M   2434M
50078   emacs        0.0  00:06.70  1    0    20   35   52K   216K  88K   18M   2392M

```

■ Running program “top” on Mac

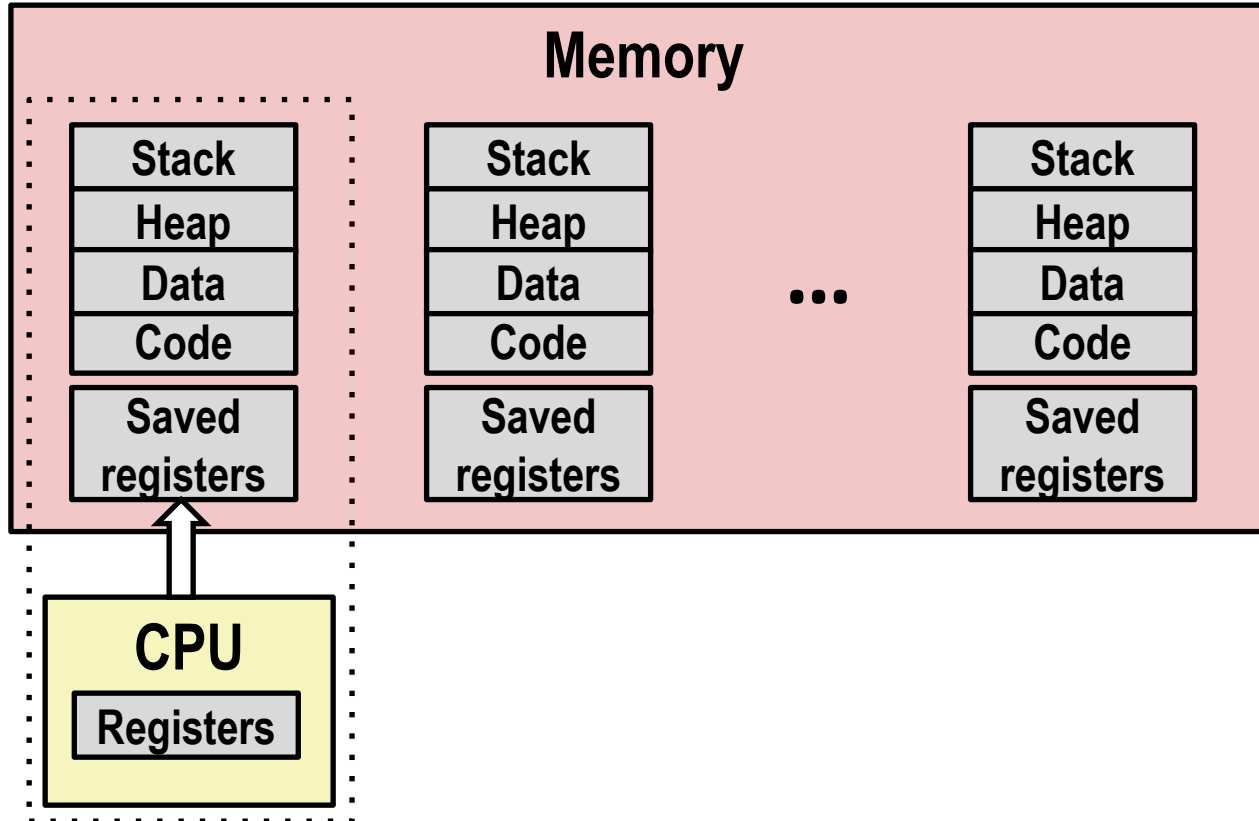
- System has 123 processes, 5 of which are active
- Identified by Process ID (PID)

Multiprocessing: Single CPU



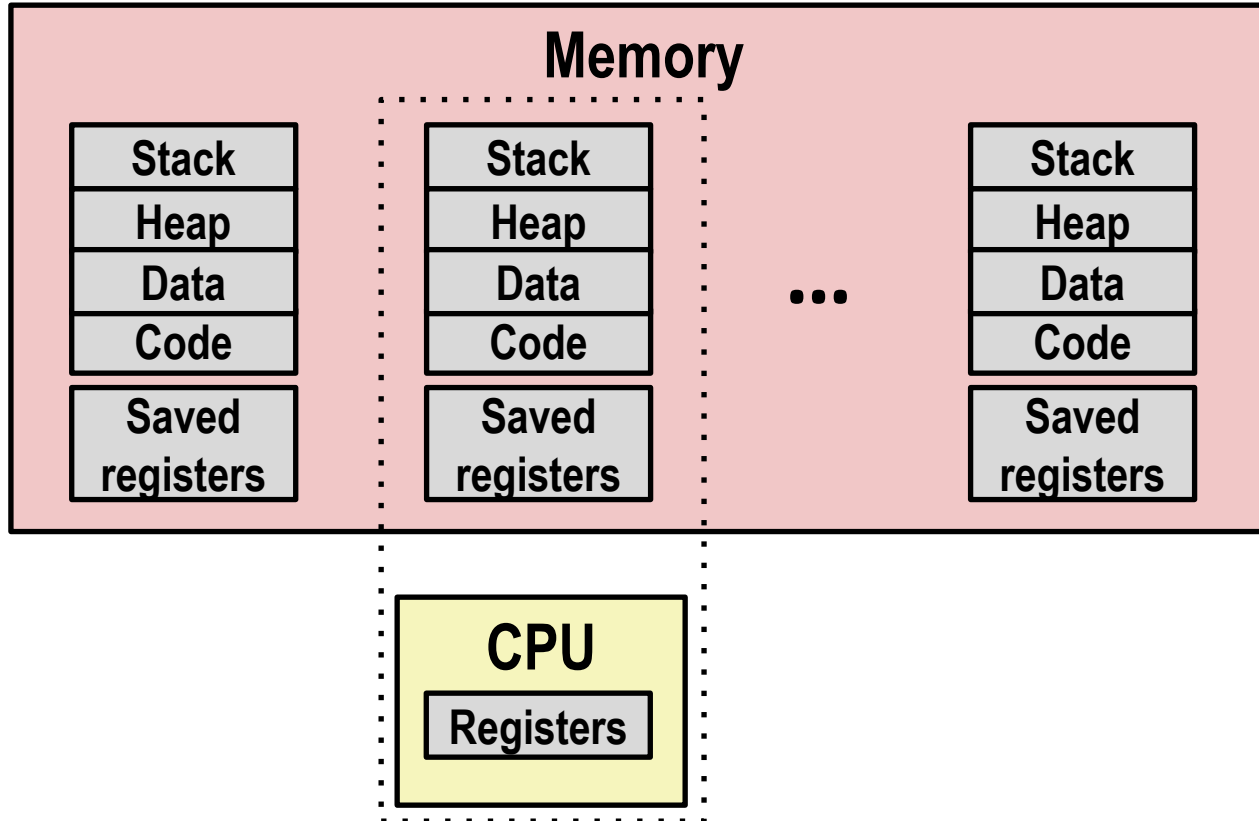
- **Single processor executes multiple processes concurrently**
 - Process executions interleaved (multitasking)
 - Address spaces managed by virtual memory system
 - Register values for nonexecuting processes saved in memory

Multiprocessing: Single CPU



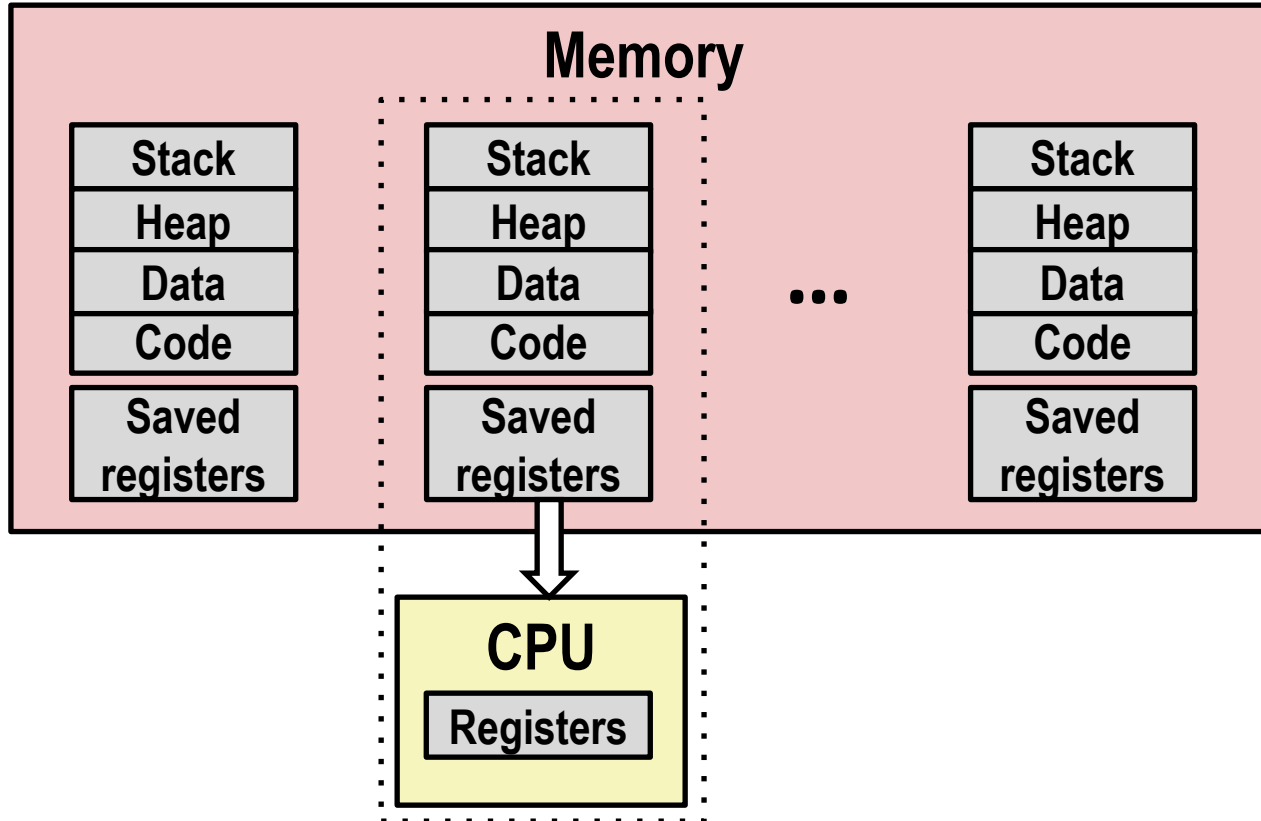
- **Save current registers in memory**

Multiprocessing: Single CPU




- **Schedule next process for execution**

Multiprocessing: Single CPU



- Load saved registers and switch address space (context switch)

Multiprocessing: Multiple CPUs (Core)

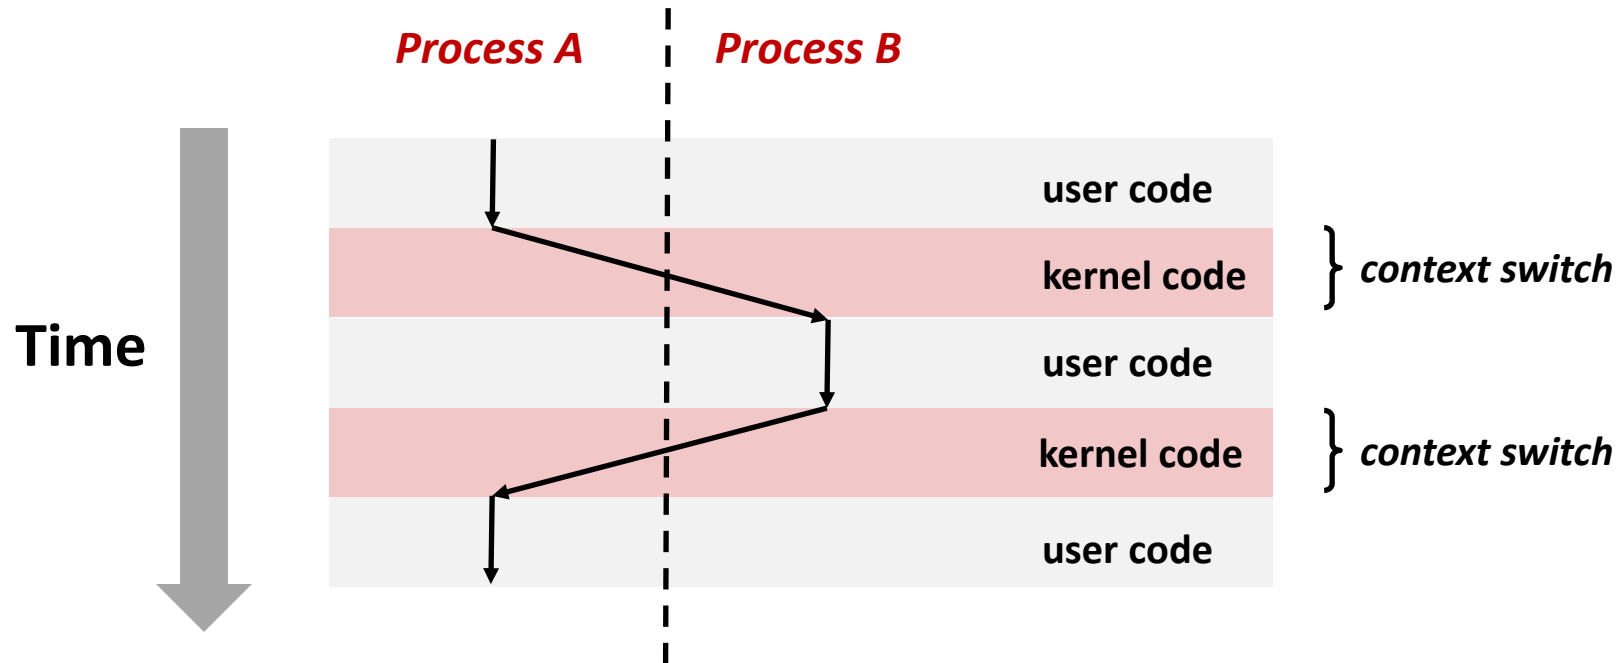


■ Multicore processors

- Multiple CPUs on single chip
- Share main memory (and some of the caches)
- Each can execute a separate process
 - Scheduling of processes onto cores done by kernel

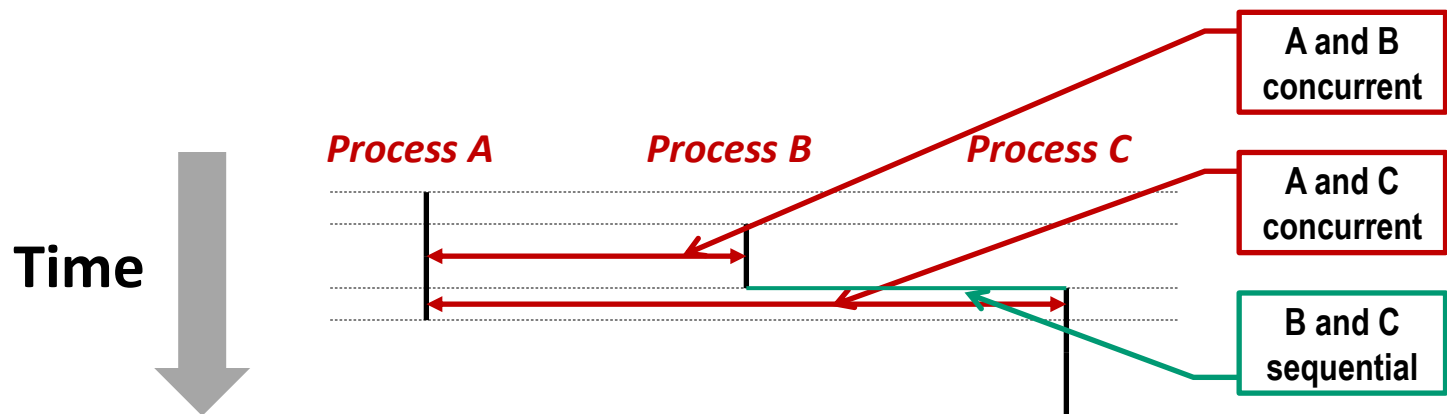
Context Switching

- Processes are managed by a shared chunk of memory-resident OS code called the *kernel*
 - Important: the kernel is not a separate process, but rather runs as part of some existing process.
- Control flow passes from one process to another via a *context switch*



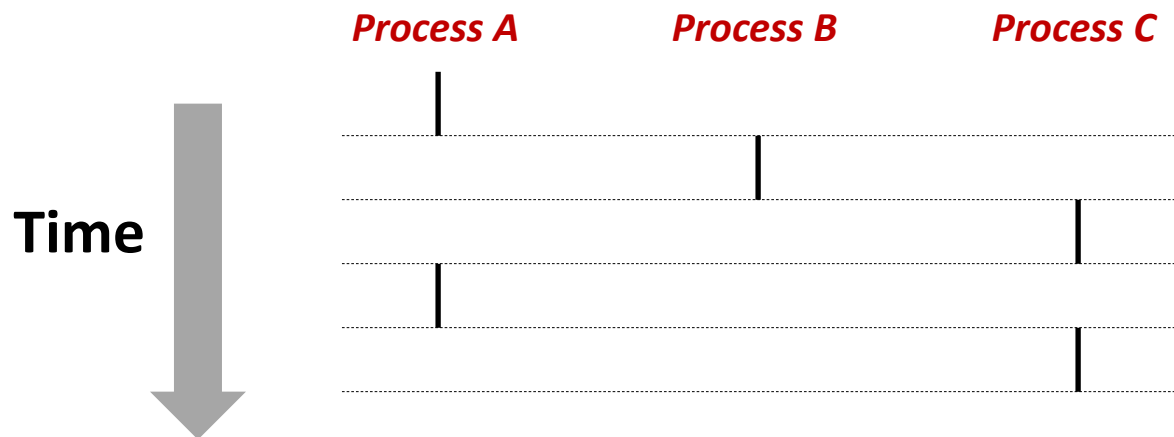
User View of Concurrent Processes

- Two processes *run concurrently* (are concurrent) if their execution overlaps in time
- Otherwise, they are *sequential*
- Appears as if concurrent processes run in parallel with each other
 - This means they can interfere with each other (e.g., synchronization) (more on that in a couple weeks)



Traditional Reality: Single CPU Case

- Only one process runs at a time
- A and B execution is *interleaved*, not truly concurrent
- Similarly for A and C
- Still possible for A and B / A and C to interfere with each other



System Calls (Revisted)

- Whenever a program wants to cause an effect outside its own process, it must ask the kernel for help
- Examples:
 - Read/write files
 - Get current time
 - Allocate RAM (sbrk)
 - Create new processes

```
// fopen.c
FILE *fopen(const char *fname,
            const char *mode) {
    int flags = mode2flags(mode);
    if (!flags) return NULL;
    int fd = open(fname, flags,
                  DEFPERMS);
    if (fd == -1) return NULL;
    return fdopen(fd, mode);
}

// open.S
.global open
open:
    mov $SYS_open, %eax
    syscall
    cmp $SYS_error_thresh, %rax
    ja  __syscall_error
    ret
```

All the System Calls

accept	fanotify_init	getresuid	llistxattr	nfsservctl	recvmsg	set_mempolicy_home_node	sync_file_range
accept4	fanotify_mark	getrlimit	lookup_dcookie	open_by_handle_at	recvmsg	set_robust_list	sync_file_range2
acct	fchdir	getrusage	lremovexattr	open_tree	remap_file_pages	set_tid_address	syncfs
add_key	fchmod	getsid	lsetxattr	openat	removexattr	setdomainname	sysinfo
adjtimex	fchmodat	getsockname	madvise	openat2	renameat	setfsuid	syslog
bind	fchown	getsockopt	mbind	perf_event_open	renameat2	setfsuid	tee
bpf	fchownat	gettid	membarrier	personality	request_key	setgid	tgkill
brk	fdatasync	gettimeofday	memfd_create	pidfd_getfd	restart_syscall	setgroups	timer_create
capget	fgetxattr	getuid	memfd_secret	pidfd_open	rseq	sethostname	timer_delete
capset	finit_module	getxattr	migrate_pages	pidfd_send_signal	rt_sigaction	setitimer	timer_getoverrun
chdir	flistxattr	init_module	mincore	pipe2	rt_sigpending	setns	timer_gettime
chroot	flock	notify_add_watch	mknodat	pivot_root	rt_sigprocmask	setpgid	timer_settime
clock_adjtime	fremovexattr	notify_init1	mknodat	pkey_alloc	rt_sigqueueinfo	setpriority	timerfd_create
clock_getres	fsconfig	notify_rm_watch	mlock	pkey_free	rt_sigreturn	setregid	timerfd_gettime
clock_gettime	fsetxattr	io_cancel	mlock2	pkey_mprotect	rt_sigsuspend	setresgid	timerfd_settime
clock_nanosleep	fsmount	io_destroy	mlockall	ppoll	rt_sigtimedwait	setresuid	times
clock_settime	fsopen	io_getevents	mount	prctl	rt_tgsigqueueinfo	setreuid	tkill
clone	fspick	io_pgetevents	mount_setattr	pread64	sched_get_priority_max	setrlimit	umask
clone3	fsync	io_setup	move_mount	preadv	sched_get_priority_min	setsid	umount2
close	futex	io_submit	move_pages	preadv2	sched_getaffinity	setsockopt	uname
close_range	futex_waitv	io_uring_enter	mprotect	prlimit64	sched_getattr	settimeofday	unlinkat
connect	get_mempolicy	io_uring_register	mq_getsetattr	process_madvise	sched_getparam	setuid	unshare
copy_file_range	get_robust_list	io_uring_setup	mq_notify	process_mrelease	sched_getscheduler	setxattr	userfaultfd
delete_module	getcpu	ioctl	mq_open	process_vm_readv	sched_rr_get_interval	shmat	utimensat
dup	getcwd	ioprio_get	mq_timedreceive	process_vm_writev	sched_setaffinity	shmctl	vhangup
dup3	getdents64	ioprio_set	mq_timedsend	pselect6	sched_setattr	shmdt	vmsplice
epoll_create1	getegid	kcmp	mq_unlink	ptrace	sched_setparam	shmget	wait4
epoll_ctl	geteuid	kexec_file_load	mremap	pwrite64	sched_setscheduler	shutdown	waitid
epoll_pwait	getgid	kexec_load	msgctl	pwritev	sched_yield	sigaltstack	write
epoll_pwait2	getgroups	keyctl	msgget	process_vm_readv	seccomp	signalfd4	writev
eventfd2	getitimer	kill	msgrcv	quotactl	semctl	socket	
execve	getpeername	landlock_add_rule	msgsnd	quotactl_fd	semget	socketpair	
execveat	getpgid	landlock_create_ruleset	msync	read	semop	splice	
exit	getpid	landlock_restrict_self	munlock	readahead	semtimedop	statx	
exit_group	getppid	lgetxattr	munlockall	readlinkat	sendmsg	swapoff	
faccessat	getpriority	linkat	munmap	readv	sendmsg	swapon	
faccessat2	getrandom	listen	name_to_handle_at	reboot	sendto	symlinkat	
fallocate	getresgid	listxattr	nanosleep	recvfrom	set_mempolicy	sync	

System Call Error Handling

- Almost all system-level operations can fail
 - You must explicitly check for failure
- On error, most system-level functions return **-1** and set global variable **errno** to indicate cause.
- Example:

```
pid_t pid = fork();  
if (pid == -1) {  
    fprintf(stderr, "fork error: %s\n", strerror(errno));  
    exit(1);  
}
```


Error-reporting Functions

- Can simplify somewhat using an *error-reporting function*:

```
void unix_error(char *msg) /* Unix-style error */
{
    fprintf(stderr, "%s: %s\n", msg, strerror(errno));
    exit(1);
}
```

```
pid_t pid = fork();
if (pid == -1)
    unix_error("fork error");
```

- Not always appropriate to exit when something goes wrong.

Error-handling Wrappers

- We simplify the code we present to you even further by using Stevens¹-style error-handling wrappers:

```
pid_t Fork(void)
{
    pid_t pid = fork();

    if (pid == -1)
        unix_error("Fork error");
    return pid;
}
```

```
pid = Fork(); // Only returns if successful
```

- NOT what you generally want to do in a real application

¹e.g., in "UNIX Network Programming: The sockets networking API" W. Richard Stevens

Obtaining Process IDs

- `pid_t getpid(void)`
 - Returns PID of current process

- `pid_t getppid(void)`
 - Returns PID of parent process

Process States

From a programmer's perspective, we can think of a process as being in one of three states

■ Running

- Process is either executing, or waiting to be executed and will eventually be *scheduled* (i.e., chosen to execute) by the kernel

■ Stopped

- Process execution is *suspended* and will not be scheduled until further notice (next lecture when we study signals)

■ Terminated

- Process is stopped permanently

Terminating Processes

■ Process becomes terminated for one of three reasons:

- Receiving a signal whose default action is to terminate (next lecture)
- Returning from the `main` routine
- Calling the `exit` function

■ `void exit(int status)`

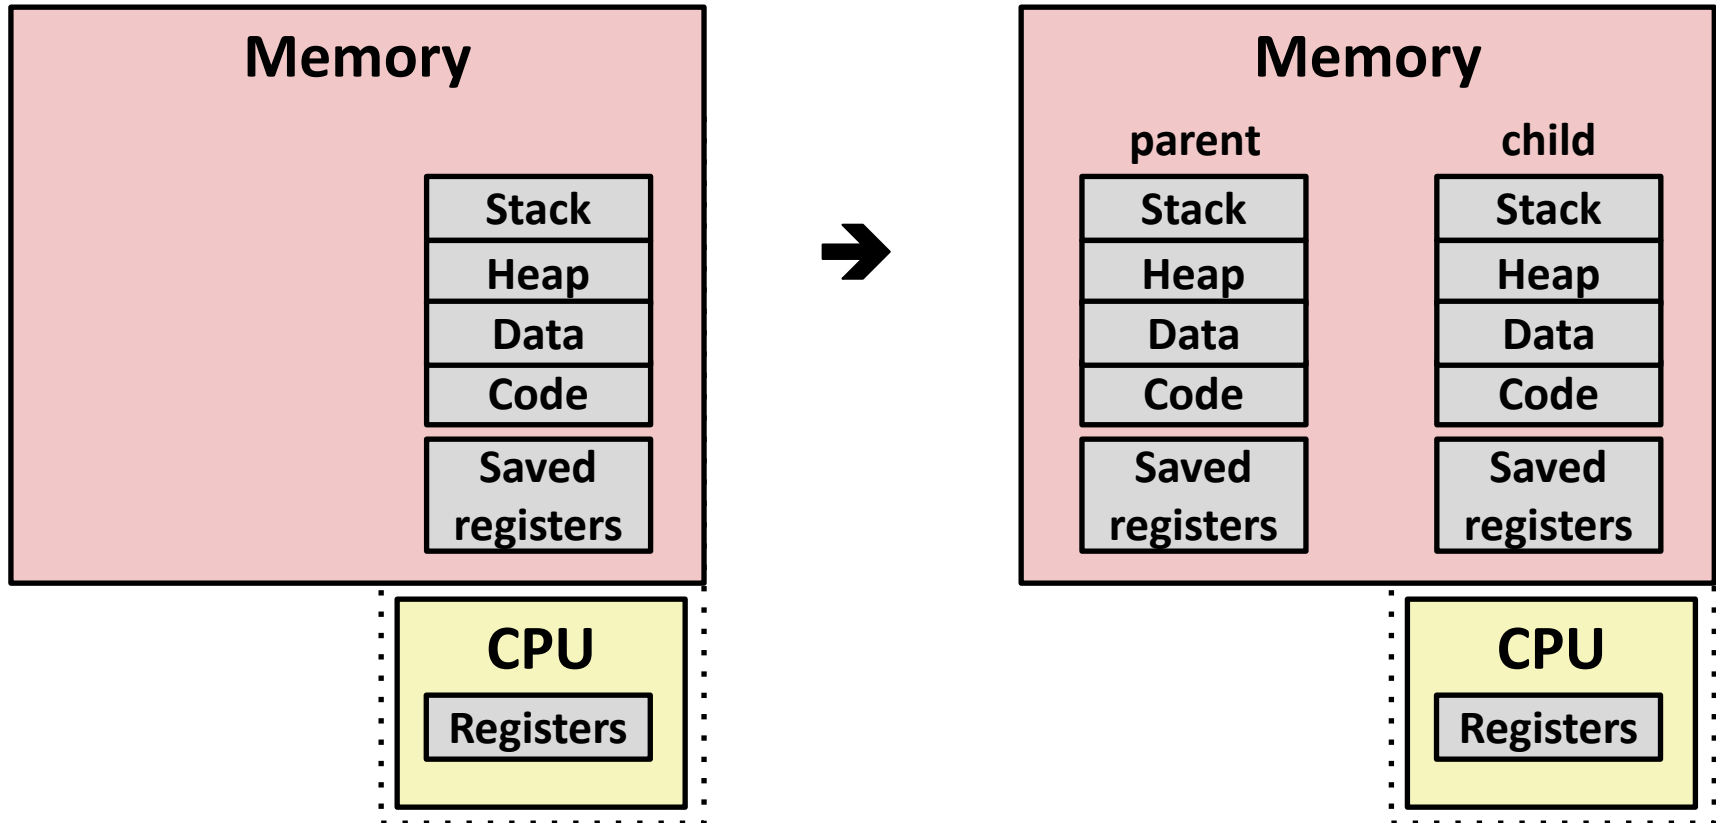
- Terminates with an *exit status* of `status`
- Convention: normal return status is 0, nonzero on error
- Another way to explicitly set the exit status is to return an integer value from the main routine

■ `exit` is called **once** but **never** returns.

Creating Processes

- *Parent process* creates a new running *child process* by calling `fork`
- `int fork(void)`
 - Returns 0 to the child process, child's PID to parent process
 - Child is *almost* identical to parent:
 - Child get an identical (but separate) copy of the parent's virtual address space.
 - Child gets identical copies of the parent's open file descriptors
 - Child has a different PID than the parent
- `fork` is interesting (and often confusing) because it is called *once* but returns *twice*

Conceptual View of fork



■ Make complete copy of execution state

- Designate one as parent and one as child
- Resume execution of parent or child
- (Optimization: Use copy-on-write to avoid copying RAM)

fork Example

```
int main(int argc, char** argv)
{
    pid_t pid;
    int x = 1;

    pid = Fork();
    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        return 0;
    }

    /* Parent */
    printf("parent: x=%d\n", --x);
    return 0;
}
```

fork.c

- Call once, return twice
- Concurrent execution
 - Can't predict execution order of parent and child

```
linux> ./fork
parent: x=0
child : x=2
```

```
linux> ./fork
child : x=2
parent: x=0
```

```
linux> ./fork
parent: x=0
child : x=2
```

```
linux> ./fork
parent: x=0
child : x=2
```


fork Example

```
int main(int argc, char** argv)
{
    pid_t pid;
    int x = 1;

    pid = Fork();
    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        return 0;
    }

    /* Parent */
    printf("parent: x=%d\n", --x);
    return 0;
}
```

```
linux> ./fork
parent: x=0
child : x=2
```

- Call once, return twice
- Concurrent execution
 - Can't predict execution order of parent and child
- Duplicate but separate address space
 - `x` has a value of 1 when fork returns in parent and child
 - Subsequent changes to `x` are independent
- Shared open files
 - `stdout` is the same in both parent and child

Modeling fork with Process Graphs

- **A *process graph* is a useful tool for capturing the partial ordering of statements in a concurrent program:**
 - Each vertex is the execution of a statement
 - $a \rightarrow b$ means a happens before b
 - Edges can be labeled with current value of variables
 - `printf` vertices can be labeled with output
 - Each graph begins with a vertex with no inedges
- **Any *topological sort* of the graph corresponds to a feasible total ordering.**
 - Total ordering of vertices where all edges point from left to right

Process Graph Example

```

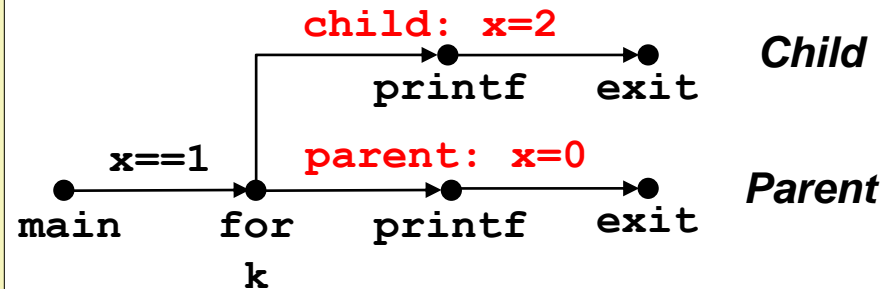
int main()
{
    pid_t pid;
    int x = 1;

    pid = Fork();
    if (pid == 0) { /* Child */
        printf("child : x=%d\n", ++x);
        exit(0);
    }

    /* Parent */
    printf("parent: x=%d\n", --x);
    exit(0);
}

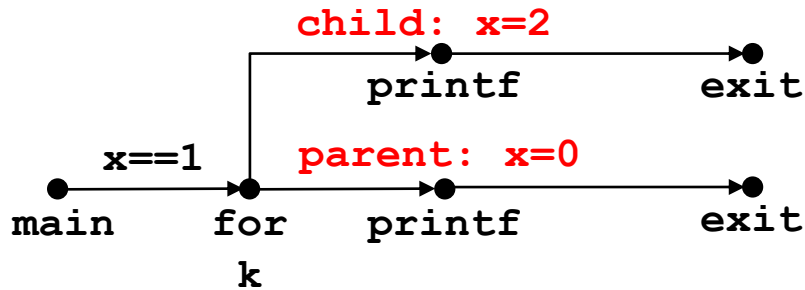
```

fork.c




Interpreting Process Graphs


■ Original graph:




■ Relabelled graph:



Feasible total ordering:



Feasible or Infeasible?




Infeasible: not a topological sort

fork Example: Two consecutive forks

```
void fork2()
{
    printf("L0\n");
    fork();
    printf("L1\n");
    fork();
    printf("Bye\n");
}
```

forks.c



Feasible output:

L0
L1
Bye
Bye
L1
Bye
Bye

Infeasible output:

L0
Bye
L1
Bye
L1
Bye
Bye

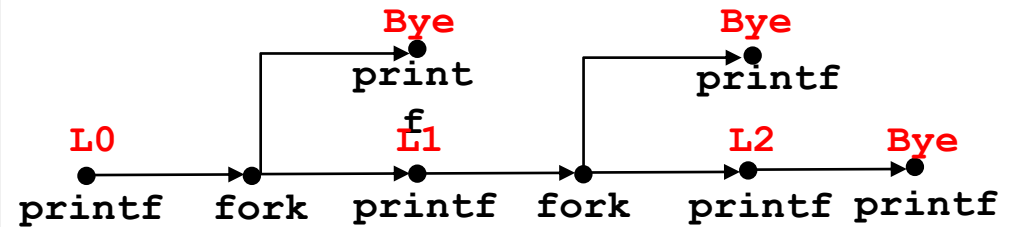
fork Example: Nested forks in parent

```

void fork4()
{
    printf("L0\n");
    if (fork() != 0) {
        printf("L1\n");
        if (fork() != 0) {
            printf("L2\n");
        }
    }
    printf("Bye\n");
}

```

forks.c



Feasible output:

L0
L1
Bye
Bye
L2
Bye

Infeasible output:

L0
Bye
L1
Bye
Bye
L2


fork Example: Nested forks in children

```

void fork5()
{
    printf("L0\n");
    if (fork() == 0) {
        printf("L1\n");
        if (fork() == 0) {
            printf("L2\n");
        }
    }
    printf("Bye\n");
}

```

forks.c



Feasible output:

L0
 Bye
 L1
 L2
 Bye
 Bye

Infeasible output:

L0
 Bye
 L1
 Bye
 Bye
 L2

Reaping Child Processes

■ Idea

- When process terminates, it still consumes system resources
 - Examples: Exit status, various OS tables
- Called a “zombie”
 - Living corpse, half alive and half dead

■ Reaping

- Performed by parent on terminated child (using `wait` or `waitpid`)
- Parent is given exit status information
- Kernel then deletes zombie child process

■ What if parent doesn't reap?

- If any parent terminates without reaping a child, then the orphaned child will be reaped by `init` process (`pid == 1`)
- So, only need explicit reaping in long-running processes
 - e.g., shells and servers

Zombie Example

```
void fork7() {
    if (fork() == 0) {
        /* Child */
        printf("Terminating Child, PID = %d\n", getpid());
        exit(0);
    } else {
        printf("Running Parent, PID = %d\n", getpid());
        while (1)
            ; /* Infinite loop */
    }
}
```

forks.c

```
linux> ./forks 7 &
[1] 6639
```

```
Running Parent, PID = 6639
```

```
Terminating Child, PID = 6640
```

```
linux> ps
```

PID	TTY	TIME	CMD
6585	ttyp9	00:00:00	tcsh
6639	ttyp9	00:00:03	forks
6640	ttyp9	00:00:00	forks <defunct>
6641	ttyp9	00:00:00	ps

```
linux> kill 6639
```

```
[1] Terminated
```

```
linux> ps
```

PID	TTY	TIME	CMD
6585	ttyp9	00:00:00	tcsh
6642	ttyp9	00:00:00	ps

■ **ps** shows child process as “defunct” (i.e., a zombie)

■ Killing parent allows child to be reaped by **init**

Non-terminating Child Example

```
void fork8()
{
    if (fork() == 0) {
        /* Child */
        printf("Running Child, PID = %d\n",
            getpid());
        while (1)
            ; /* Infinite loop */
    } else {
        printf("Terminating Parent, PID = %d\n",
            getpid());
        exit(0);
    }
}
```

forks.c

```
linux> ./forks 8
Terminating Parent, PID = 6675
Running Child, PID = 6676
linux> ps
  PID TTY          TIME CMD
 6585 tttyp9      00:00:00 tcsh
 6676 tttyp9      00:00:06 forks
 6677 tttyp9      00:00:00 ps
linux> kill 6676
linux> ps
  PID TTY          TIME CMD
 6585 tttyp9      00:00:00 tcsh
 6678 tttyp9      00:00:00 ps
```

■ Child process still active even though parent has terminated

■ Must kill child explicitly, or else will keep running indefinitely

`wait`: Synchronizing with Children

- Parent reaps a child by calling the `wait` function
- `int wait(int *child_status)`
 - Suspends current process until one of its children terminates
 - Return value is the `pid` of the child process that terminated
 - If `child_status != NULL`, then the integer it points to will be set to a value that indicates reason the child terminated and the exit status:
 - Checked using macros defined in `wait.h`
 - `WIFEXITED`, `WEXITSTATUS`, `WIFSIGNALED`,
`WTERMSIG`, `WIFSTOPPED`, `WSTOPSIG`,
`WIFCONTINUED`
 - See textbook for details

wait: Synchronizing with Children


```

void fork9() {
    int child_status;

    if (fork() == 0) {
        printf("HC: hello from child\n");
        exit(0);
    } else {
        printf("HP: hello from parent\n");
        wait(&child_status);
        printf("CT: child has terminated\n");
    }
    printf("Bye\n");
}

```

forks.c



Feasible output:

HC
HP
CT
Bye

Infeasible output:

HP
CT
Bye
HC

Another wait Example

- If multiple children completed, will take in arbitrary order
- Can use macros WIFEXITED and WEXITSTATUS to get information about exit status

```
void fork10() {  
    pid_t pid[N];  
    int i, child_status;  
  
    for (i = 0; i < N; i++)  
        if ((pid[i] = fork()) == 0) {  
            exit(100+i); /* Child */  
        }  
    for (i = 0; i < N; i++) { /* Parent */  
        pid_t wpid = wait(&child_status);  
        if (WIFEXITED(child_status))  
            printf("Child %d terminated with exit status %d\n",  
                wpid, WEXITSTATUS(child_status));  
        else  
            printf("Child %d terminate abnormally\n", wpid);  
    }  
}
```

forks.c

waitpid: Waiting for a Specific Process

- `pid_t waitpid(pid_t pid, int &status, int options)`
 - Suspends current process until specific process terminates
 - Various options (see textbook)

```
void fork11() {  
    pid_t pid[N];  
    int i;  
    int child_status;  
  
    for (i = 0; i < N; i++)  
        if ((pid[i] = fork()) == 0)  
            exit(100+i); /* Child */  
    for (i = N-1; i >= 0; i--) {  
        pid_t wpid = waitpid(pid[i], &child_status, 0);  
        if (WIFEXITED(child_status))  
            printf("Child %d terminated with exit status %d\n",  
                wpid, WEXITSTATUS(child_status));  
        else  
            printf("Child %d terminate abnormally\n", wpid);  
    }  
}
```

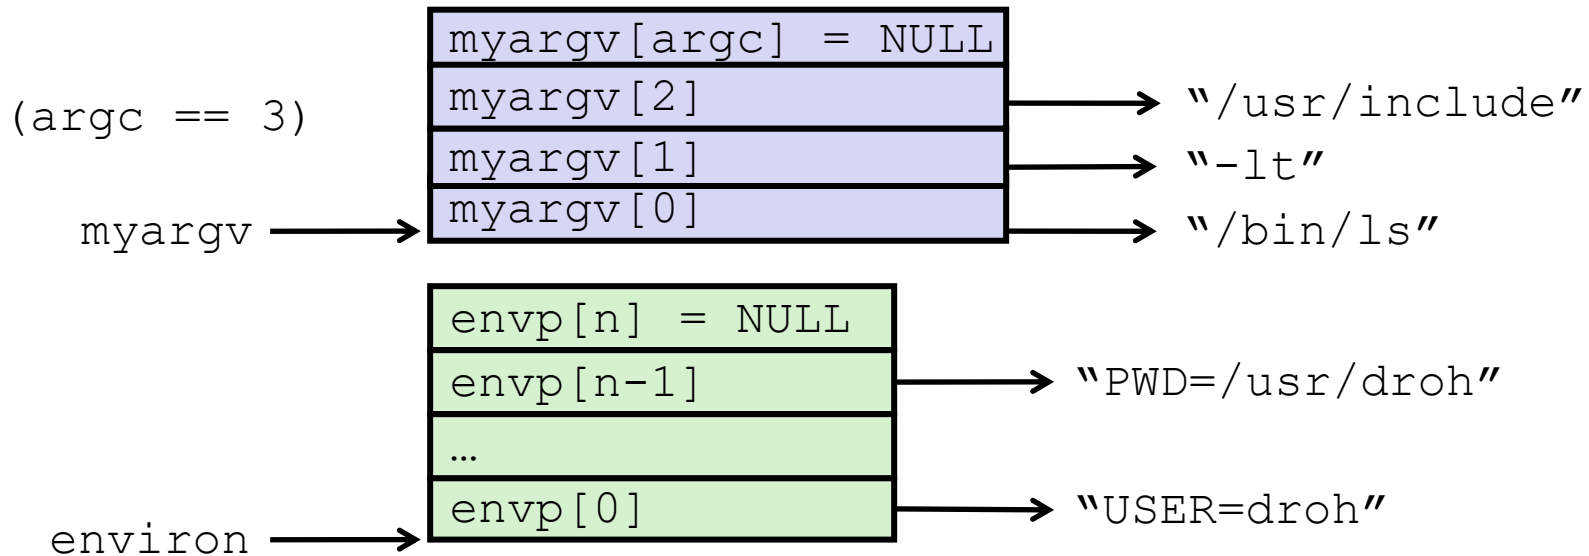
forks.c

execve : Loading and Running Programs

- `int execve(char *filename, char *argv[], char *envp[])`
- **Loads and runs in the current process:**
 - Executable file `filename`
 - Can be object file or script file beginning with `#!interpreter` (e.g., `#!/bin/bash`)
 - ...with argument list `argv`
 - By convention `argv[0]==filename`
 - ...and environment variable list `envp`
 - “name=value” strings (e.g., `USER=droh`)
 - `getenv`, `putenv`, `putenv`
- **Overwrites code, data, and stack**
 - Retains PID, open files and signal context
- **Called **once** and **never** returns**
 - ...except if there is an error

execve Example

- Executes `"/bin/ls -lt /usr/include"` in child process using current environment:



```

if ((pid = Fork()) == 0) { /* Child runs program */
    if (execve(myargv[0], myargv, environ) < 0) {
        printf("%s: Command not found.\n", myargv[0]);
        exit(1);
    }
}

```


Summary

■ Exceptions

- Events that require nonstandard control flow
- Generated externally (interrupts) or internally (traps and faults)

■ Processes

- At any given time, system has multiple active processes
- Only one can execute at a time on a single core, though
- Each process appears to have total control of processor + private memory space

Summary (cont.)

■ Spawning processes

- Call `fork`
- One call, two returns

■ Process completion

- Call `exit`
- One call, no return

■ Reaping and waiting for processes

- Call `wait` or `waitpid`

■ Loading and running programs

- Call `execve` (or variant)
- One call, (normally) no return