

Synchronization

CSE4100: Multicore Programming

Sungyong Park (PhD)

Data-Intensive Computing and Systems Laboratory (DISCOS)

<https://discos.sogang.ac.kr>

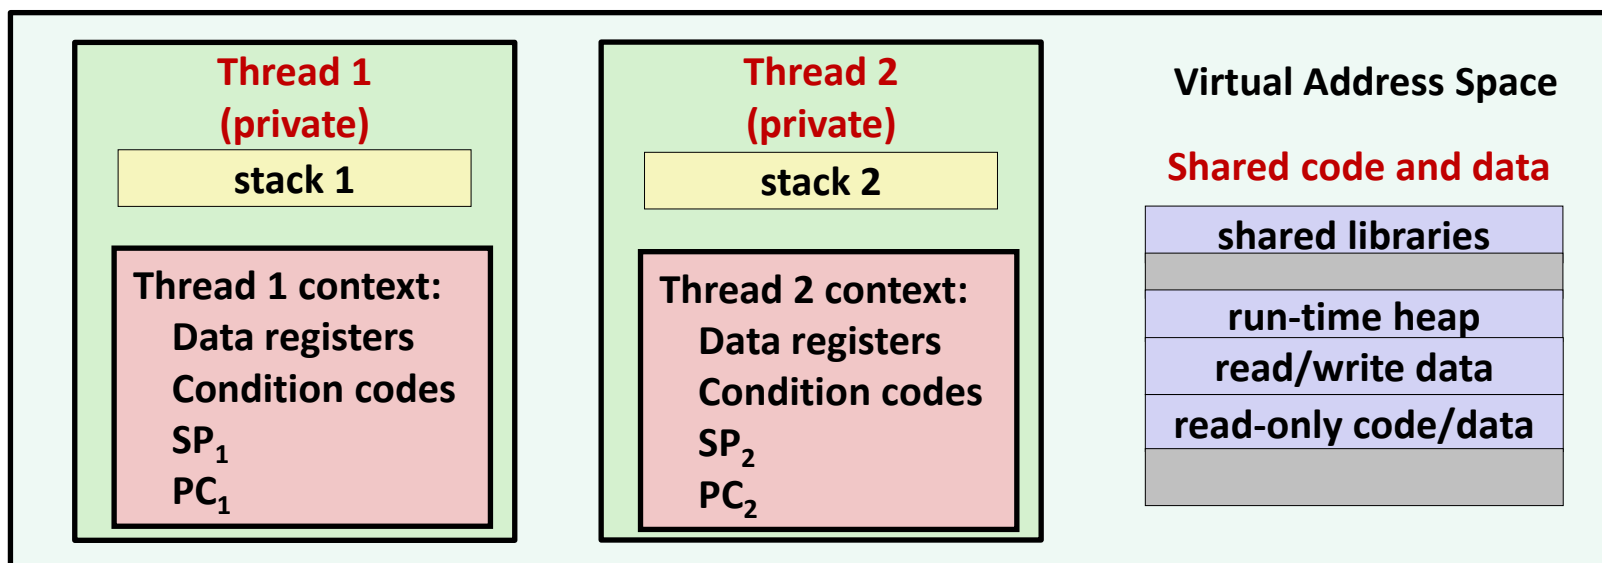
Office: R908A, E-mail: parksy@sogang.ac.kr

Shared Variables in Threaded C Programs

- **Question: Which variables in a threaded C program are shared?**
 - The answer is not as simple as “*global variables are shared*” and “*stack variables are private*”
- **A variable x is *shared* if and only if (iff) multiple threads reference some instance of x**
- **Requires answers to the following questions**
 - What is the memory model for threads?
 - How are instances of variables mapped to memory?
 - How many threads might reference each of these instances?

Threads Memory Model: Conceptual

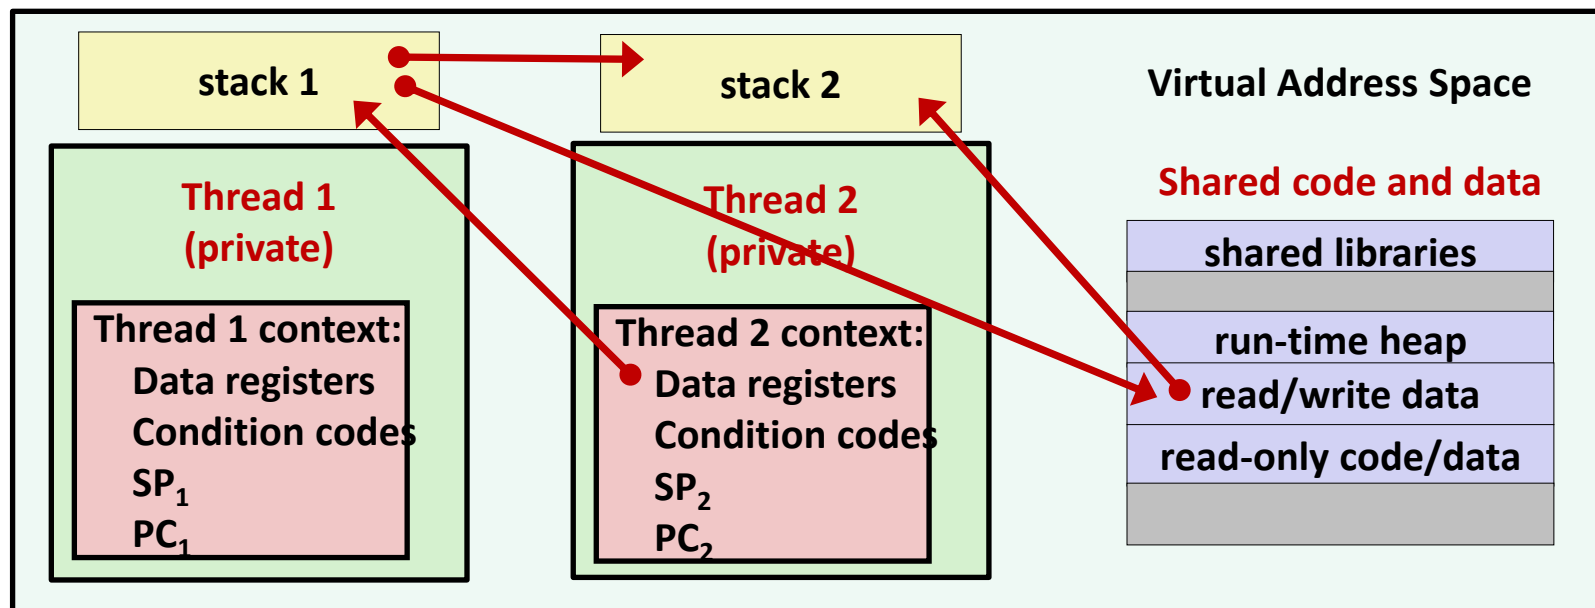
- Multiple threads run within the context of a single process
- Each thread has its own separate thread context
 - Thread ID, stack, stack pointer, PC, condition codes, and GP registers
- All threads share the remaining process context
 - Code, data, heap, and shared library segments of the process virtual address space
 - Open files and installed handlers



Threads Memory Model: Actual

■ Separation of data is not strictly enforced:

- Register values are truly separate and protected, but...
- Any thread can read and write the stack of any other thread



The mismatch between the conceptual and operation model is a source of confusion and errors!

Three Ways to Pass Thread Argument

■ Malloc / Free

- Producer malloc's space, passes pointer to `pthread_create`
- Consumer dereferences pointer, frees space
- Always works; necessary for passing large amounts of data

■ Cast of int

- Producer casts an int/long to `void*`, passes to `pthread_create`
- Consumer casts `void*` argument back to int/long
- Works for small amounts of data (one number)

■ **INCORRECT: Pointer to stack slot**

- Producer passes address to producer's stack in `pthread_create`
- Consumer dereferences pointer
- Why is this unsafe?

Passing an Argument to a Thread (Case 1)

```
int hist[N] = {0};

int main(int argc, char *argv[])
{
    long i;
    pthread_t tids[N];

    for ( i = 0; i < N; i++ )
        Pthread_create(&tids[i], NULL,
                      thread, &hist[i]);

    for ( i = 0; i < N; i++ )
        Pthread_join(tids[i], NULL);
    check();
}
```

```
void *thread(void *vargp)
{
    *(int *)vargp += 1;
    return NULL;
}
```

```
void check(void) {
    for ( int i = 0; i < N; i++ ) {
        if ( hist[i] != 1 ) {
            printf("Failed at %d\n", i);
            exit(-1);
        }
    }
    printf("OK\n");
}
```

- Each thread receives a *unique pointer*

Passing an Argument to a Thread (Case 2)

```
int hist[N] = {0};

int main(int argc, char *argv[]) {
    long i;
    pthread_t tids[N];

    for (i = 0; i < N; i++)
        Pthread_create(&tids[i],
                       NULL,
                       thread,
                       (void *)i);
    for (i = 0; i < N; i++)
        Pthread_join(tids[i], NULL);
    check();
}
```

```
void *thread(void *vargp)
{
    hist[(long)vargp] += 1;
    return NULL;
}
```

- Each thread receives a *unique array index*
- Casting from long to void* and back is safe

Passing an Argument to a Thread (Case 3)

```
int hist[N] = {0};

int main(int argc, char *argv[]) {
    long i;
    pthread_t tids[N];

    for ( i = 0; i < N; i++ )
        long* p = Malloc(sizeof(long));
        *p = i;
        Pthread_create(&tids[i], NULL,
                      thread, p);

    for ( i = 0; i < N; i++ )
        Pthread_join(tids[i], NULL);
    check();
}
```

```
void *thread(void *vargp)
{
    hist[*(long *)vargp] += 1;
    free(vargp);
    return NULL;
}
```

- Each thread receives a *unique array index*
- Malloc in parent, free in thread
- Necessary if passing structs

Passing an Argument to a Thread – **WRONG!**

```
int hist[N] = {0};

int main(int argc, char *argv[]) {
    long i;
    pthread_t tids[N];

    for ( i = 0; i < N; i++ )
        Pthread_create(&tids[i], NULL,
                      thread, &i);

    for ( i = 0; i < N; i++ )
        Pthread_join(tids[i], NULL);
    check();
}
```

```
void *thread(void *vargp)
{
    hist[*(long *)vargp] += 1;
    return NULL;
}
```

- Each thread receives *the same pointer*, to *i* in main
- Data race: each thread *may or may not* read a unique array index from *i* in main

Mapping Variable Instances to Memory

■ Global variables

- Variable declared outside of a function
- **Virtual memory contains exactly one instance of any global variable**

■ Local automatic variables

- Variable declared inside function without `static` attribute
- **Each thread stack contains one instance of each local variable**

■ Local static variables

- Variable declared inside function with the `static` attribute
- **Virtual memory contains exactly one instance of any local static variable**

■ `errno` is special

- Declared outside a function, but **each thread stack contains one instance**

Mapping Variable Instances to Memory (Cont)

Global var: 1 instance (ptr [data])

```
char **ptr; /* global var */
```

```
int main()
```

```
{
```

```
    long i;
```

```
    pthread_t tid;
```

```
    char *msgs[2] = {  
        "Hello from foo",  
        "Hello from bar"  
    };
```

```
    ptr = msgs;
```

```
    for (i = 0; i < 2; i++)  
        Pthread_create(&tid,  
                        NULL, thread, (void *)i);  
    Pthread_exit(NULL);
```

```
}
```

sharing.c

Local var: 1 instance (i.m, msgs.m)

Local var: 2 instances (
 myid.p0 [peer thread 0's stack],
 myid.p1 [peer thread 1's stack]
)

```
void *thread(void *vargp)
```

```
{
```

```
    long myid = (long)vargp;
```

```
    static int cnt = 0;
```

```
    printf("[%ld]:  %s (cnt=%d)\n",  
           myid, ptr[myid], ++cnt);
```

```
    return NULL;
```

```
}
```

Local static var: 1 instance (cnt [data])

Shared Variable Analysis

■ Which variables are shared?

<i>Variable instance</i>	<i>Referenced by main thread?</i>	<i>Referenced by peer thread 0?</i>	<i>Referenced by peer thread 1?</i>	
ptr	yes	yes	yes	→ <i>global</i>
cnt	no	yes	yes	→ <i>static</i>
i.m	yes	no	no	
msgs.m	yes	yes	yes	→ <i>global via ptr</i>
myid.p0	no	yes	no	
myid.p1	no	no	yes	

- Answer: A variable **x** is shared if and only if one of its instances is referenced by more than one thread
 - **ptr**, **cnt**, and **msgs** are shared
 - **i** and **myid** are *not* shared

Synchronizing Threads

- Shared variables are handy...
- ...but introduce the possibility of nasty *synchronization* errors

badcnt.c: Improper Synchronization

```

/* Global shared variable */
volatile long cnt = 0; /* Counter */

int main(int argc, char **argv)
{
    long niters;
    pthread_t tid1, tid2;

    niters = atoi(argv[1]);
    Pthread_create(&tid1, NULL,
        thread, &niters);
    Pthread_create(&tid2, NULL,
        thread, &niters);
    Pthread_join(tid1, NULL);
    Pthread_join(tid2, NULL);

    /* Check result */
    if (cnt != (2 * niters))
        printf("BOOM! cnt=%ld\n", cnt);
    else
        printf("OK cnt=%ld\n", cnt);
    exit(0);
}

```

badcnt.c

```

/* Thread routine */
void *thread(void *vargp)
{
    long i;
    long niters = *((long *)vargp);

    for (i = 0; i < niters; i++)
        cnt++;

    return NULL;
}

```

```

linux> ./badcnt 10000
OK cnt=20000
linux> ./badcnt 10000
BOOM! cnt=13051
linux>

```

cnt should be equal to 20,000!

What went wrong?

Assembly Code for Counter Loop

C code for counter loop in thread i

```
for ( i = 0; i < niters; i++ )
    cnt++;
```

Assembly code for thread i

<pre> movq (%rdi), %rcx testq %rcx,%rcx jle .L2 movl \$0, %eax </pre>	} H_i : Head
<pre> .L3: movq cnt(%rip), %rdx addq \$1, %rdx movq %rdx, cnt(%rip) </pre>	L_i : Load cnt U_i : Update cnt S_i : Store cnt
<pre> addq \$1, %rax cmpq %rcx, %rax jne .L3 .L2: </pre>	} T_i : Tail

Concurrent Execution (Case 1)

- **Key idea:** In general, any sequentially consistent interleaving is possible, but some give an unexpected result!

- I_i denotes that thread i executes instruction I
- $\%rdx_i$ is the content of $\%rdx$ in thread i 's context

Thread 1 **critical section**
 Thread 2 **critical section**

i (thread)	$instr_i$	$\%rdx_1$	$\%rdx_2$	cnt
1	H_1	-	-	0
1	L_1	0	-	0
1	U_1	1	-	0
1	S_1	1	-	1
2	H_2	-	-	1
2	L_2	-	1	1
2	U_2	-	2	1
2	S_2	-	2	2
2	T_2	-	2	2
1	T_1	1	-	2

OK

Concurrent Execution (Case 2)

- Incorrect ordering: two threads increment the counter, but the result is 1 instead of 2

i (thread)	instr_i	$\%rdx_1$	$\%rdx_2$	cnt
1	H_1	-	-	0
1	L_1	0	-	0
1	U_1	1	-	0
2	H_2	-	-	0
2	L_2	-	0	0
1	S_1	1	-	1
1	T_1	1	-	1
2	U_2	-	1	1
2	S_2	-	1	1
2	T_2	-	1	1

Oops!

Concurrent Execution (Case 3)

■ How about this ordering?


i (thread)	instr_i	$\%rdx_1$	$\%rdx_2$	cnt
1	H_1			0
1	L_1	0		
2	H_2			
2	L_2		0	
2	U_2		1	
2	S_2		1	1
1	U_1	1		
1	S_1	1		1
1	T_1			1
2	T_2			1

Oops!

■ We can analyze the behavior using a *progress graph*

Progress Graphs


Thread 2



- A **progress graph** depicts the discrete **execution state space** of concurrent threads
- Each axis corresponds to the sequential order of instructions in a thread
- Each point corresponds to a possible **execution state** ($Inst_1$, $Inst_2$)
- For example, (L_1, S_2) denotes state where thread 1 has completed L_1 and thread 2 has completed S_2


Trajectories in Progress Graphs

Thread 2




- A **trajectory** is a sequence of legal state transitions that describes one possible concurrent execution of the threads
- For example, $H_1, L_1, U_1, H_2, L_2, S_1, T_1, U_2, S_2, T_2$ is a trajectory

Critical Sections and Unsafe Regions



- L , U , and S form a **critical section** with respect to the shared variable `cnt`
- Instructions in critical sections (w.r.t some shared variable) should not be interleaved
- Sets of states where such interleaving occurs form **unsafe regions**

Critical Sections and Unsafe Regions (Cont)



- **Def:** A trajectory is *safe* iff it does not enter any unsafe region
- **Claim:** A trajectory is correct (w.r.t cnt) iff it is safe

Enforcing Mutual Exclusion

- **Question:** How can we guarantee a safe trajectory?
- **Answer:** We must *synchronize* the execution of the threads so that they can never have an unsafe trajectory
 - i.e., need to guarantee *mutually exclusive access* for each critical section
- **Classic solution:**
 - Semaphores (Edsger Dijkstra)
- **Other approaches (out of our scope)**
 - Mutex and condition variables (Pthreads)
 - Monitors (Java)

Semaphores

- ***Semaphore***: non-negative global integer synchronization variable manipulated by ***P*** and ***V*** operations
- ***P(s) : wait(s)***
 - If ***s*** is nonzero, then decrement ***s*** by 1 and return immediately
 - Test and decrement operations occur atomically (indivisibly)
 - If ***s*** is zero, then suspend the thread until ***s*** becomes nonzero and the thread is restarted by a ***V*** operation
 - After restarting, the ***P*** operation decrements ***s*** and returns control to the caller
- ***V(s) : signal(s)***
 - Increment ***s*** by 1 (Increment operation occurs atomically)
 - If there are any threads blocked in a ***P*** operation waiting for ***s*** to become non-zero, then restart *exactly one of those threads* (order or execution is implementation dependent but possibly the 1st thread on the queue), which then completes its ***P*** operation by decrementing ***s***
- **Semaphore invariant: ($s \geq 0$)**

Semaphore Operations in C

■ Pthreads functions

```
#include <semaphore.h>

int sem_init(sem_t *s, 0, unsigned int val); /* s = val */

int sem_wait(sem_t *s); /* P(s) */
int sem_post(sem_t *s); /* V(s) */
```

■ CS:APP wrapper functions

```
#include "csapp.h"

void P(sem_t *s); /* Wrapper function for sem_wait */
void V(sem_t *s); /* Wrapper function for sem_post */
```

Fix Improper Synchronization (badcnt.c)

```

/* Global shared variable */
volatile long cnt = 0; /* Counter */

int main(int argc, char **argv)
{
    long niters;
    pthread_t tid1, tid2;

    niters = atoi(argv[1]);
    Pthread_create(&tid1, NULL,
        thread, &niters);
    Pthread_create(&tid2, NULL,
        thread, &niters);
    Pthread_join(tid1, NULL);
    Pthread_join(tid2, NULL);

    /* Check result */
    if (cnt != (2 * niters))
        printf("BOOM! cnt=%ld\n", cnt);
    else
        printf("OK cnt=%ld\n", cnt);
    exit(0);
}

```

badcnt.c

```

/* Thread routine */
void *thread(void *vargp)
{
    long i;
    long niters = *((long *)vargp);

    for (i = 0; i < niters; i++)
        cnt++;


    return NULL;
}

```

```

linux> ./badcnt 10000
OK cnt=20000
linux> ./badcnt 10000
BOOM! cnt=13051
linux>

```



How can we fix this using
semaphores?

Using Semaphores for Mutual Exclusion

■ Basic idea

- Associate a unique semaphore *mutex*, initially 1, with each shared variable (or related set of shared variables)
- Surround corresponding critical section (CS) with *P(mutex)* and *V(mutex)* operation

```
mutex = 1  
  
P(mutex)  
CS  
V(mutex)
```

■ Terminology

- **Binary semaphore**: semaphore whose value is always 0 or 1
- **Mutex**: binary semaphore used for mutual exclusion
 - P operation: *locking* the mutex
 - V operation: *unlocking* or *releasing* the mutex
 - *Holding* a mutex: locked and not yet unlocked
- **Counting semaphore**: used as a counter for set of available resources

goodcnt.c: Proper Synchronization

- Define and initialize a mutex for the shared variable cnt

```
volatile long cnt = 0; /* Counter */
sem_t mutex;          /* Semaphore that protects cnt */

Sem_init(&mutex, 0, 1); /* mutex = 1 */
```

- Surround critical section with *P* and *V*

```
for (i = 0; i < niters; i++) {
    P(&mutex);
    cnt++;
    V(&mutex);
}
```


goodcnt.c

```
linux> ./goodcnt 10000
OK cnt=20000
linux> ./goodcnt 10000
OK cnt=20000
linux>
```

Warning: It is orders of magnitude slower than badcnt.c

Why Mutexes Work

Thread 2



- Provide mutually exclusive access to shared variable by surrounding critical section with P and V operations on semaphore s (initially set to 1)
- Semaphore invariant creates a **forbidden region** that encloses unsafe region and that cannot be entered by any trajectory


Using Semaphores to Coordinate Access to Shared Resources

- **Basic idea: Thread uses a semaphore operation to notify another thread that some condition has become true**
 - Use counting semaphores to keep track of resource state and to notify other threads
 - Use *mutex* to protect access to resource
- **Two classic examples**
 - The *Producer-Consumer* Problem
 - The *Readers-Writers* Problem

Producer-Consumer Problem

■ Common synchronization pattern

- Producer waits for an empty *slot*, inserts an item into the buffer, and notifies consumer
- Consumer waits for an *item*, removes it from the buffer, and notifies producer



■ Examples

- Multimedia processing:
 - Producer creates MPEG video frames, consumer renders them
- Event-driven graphical user interfaces:
 - Producer detects mouse clicks, mouse movements, and keyboard hits and inserts corresponding events in buffer
 - Consumer retrieves events from buffer and paints the display

Producer-Consumer on an n -Element Buffer

- Requires a *mutex* and two counting semaphores
 - **mutex**: enforces mutually exclusive access to the buffer
 - **slots**: counts the available slots in the buffer
 - **items**: counts the available items in the buffer
- Implemented using a shared buffer package called **sbuf**

```
#include "csapp.h"

typedef struct {
    int *buf;           /* Buffer array */
    int n;              /* Maximum number of slots */
    int front;          /* buf[(front+1)%n] is the first item */
    int rear;           /* buf[rear%n] is the last item */
    sem_t mutex;        /* Protects accesses to buf */
    sem_t slots;        /* Counts available slots */
    sem_t items;        /* Counts available items */
} sbuf_t;

void sbuf_init(sbuf_t *sp, int n);
void sbuf_deinit(sbuf_t *sp);
void sbuf_insert(sbuf_t *sp, int item);
int sbuf_remove(sbuf_t *sp);
```


sbuf Package – Initialization

■ Initializing and de-initializing a shared buffer

```
/* Create an empty, bounded, shared FIFO buffer with n slots */
void sbuf_init(sbuf_t *sp, int n)
{
    sp->buf = Calloc(n, sizeof(int));
    sp->n = n; /* Buffer holds max of n items */
    sp->front = sp->rear = 0; /* Empty buffer iff front == rear */
    Sem_init(&sp->mutex, 0, 1); /* Binary semaphore for locking */
    Sem_init(&sp->slots, 0, n); /* Initially, buf has n empty slots */
    Sem_init(&sp->items, 0, 0); /* Initially, buf has 0 items */
}

/* Clean up buffer sp */
void sbuf_deinit(sbuf_t *sp)
{
    Free(sp->buf);
}
```

sbuf.c

sbuf Package – Insert / Remove

■ Inserting / removing an item into / from a shared buffer

```

/* Insert an item onto the rear of shared buffer sp */
void sbuf_insert(sbuf_t *sp, int item)
{
    P(&sp->slots);           /* Wait for available slot */
    P(&sp->mutex);           /* Lock the buffer */
    sp->buf[ (++sp->rear) % (sp->n) ]
        = item;              /* Insert the item */
    V(&sp->mutex);           /* Unlock the buffer */
    V(&sp->items);           /* Announce available item */
}

```

```

/* Remove and return the first item from buffer sp */
int sbuf_remove(sbuf_t *sp)
{
    int item;
    P(&sp->items);           /* Wait for available item */
    P(&sp->mutex);           /* Lock the buffer */
    item =
        sp->buf[ (++sp->front) % (sp->n) ]; /* Remove the item */
    V(&sp->mutex);           /* Unlock the buffer */
    V(&sp->slots);           /* Announce available slot */
    return item;
}

```

Readers-Writers Problem

■ Generalization of the mutual exclusion problem

■ Problem statement

- *Reader* threads only read the object
- *Writer* threads modify the object
- Writers must have exclusive access to the object
- Unlimited number of readers can access the object

Can new reader/writer run?

Currently running {

	<i>Reader</i>	<i>Writer</i>
<i>Reader</i>	O	X
<i>Writer</i>	X	X

■ Occurs frequently in real systems, e.g.,

- Online airline reservation system
- Multithreaded caching Web proxy

Variants of Readers-Writers

■ ***First readers-writers problem*** (favors readers)

- No reader should be kept waiting unless a writer has already been granted permission to use the object
- A reader that arrives after a waiting writer gets priority over the writer

■ ***Second readers-writers problem*** (favors writers)

- Once a writer is ready to write, it performs its write as soon as possible
- A reader that arrives after a writer must wait, even if the writer is also waiting

■ ***Starvation*** (where a thread waits indefinitely) is possible in both cases

Solution to First Readers-Writers Problem

Readers

```

int readcnt;      /* Initially = 0 */
sem_t mutex, w;  /* Initially = 1 */

void reader(void)
{
    while ( 1 ) {
        P(&mutex);
        readcnt++;
        if ( readcnt == 1 ) /* First in */
            P(&w);
        V(&mutex);

        /* CS: Reading happens */

        P(&mutex);
        readcnt--;
        if ( readcnt == 0 ) /* Last out */
            V(&w);
        V(&mutex);
    }
}

```

Writers

```

void writer(void)
{
    while ( 1 ) {
        P(&w);

        /* CS: Writing happens */

        V(&w);
    }
}

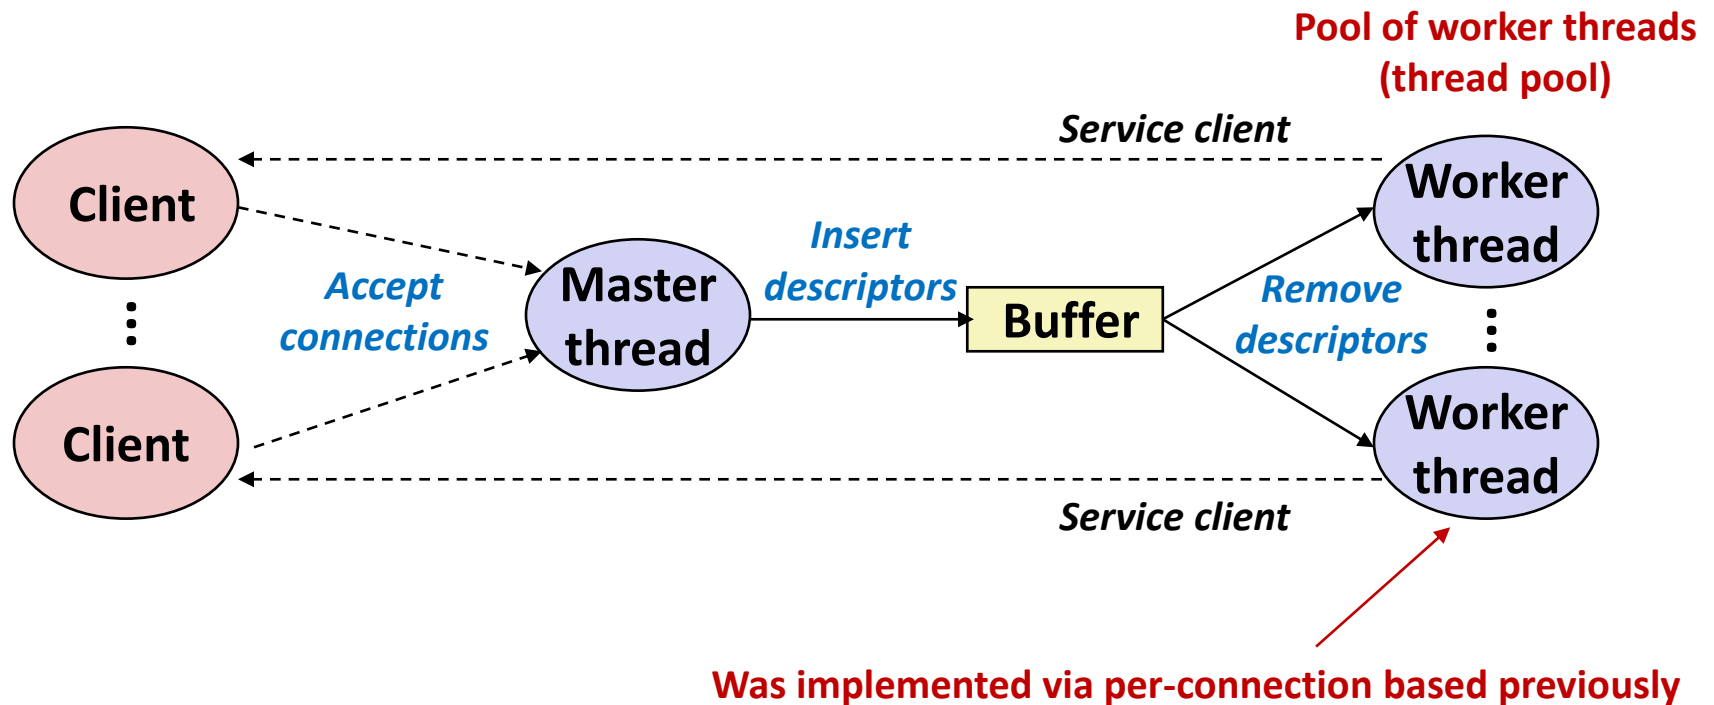
```

rw1.c

What is the order of execution for the following readers/writers?

- (1) $R_1 \rightarrow W_1 \rightarrow R_2 \rightarrow W_2$
- (2) $W_1 \rightarrow R_1 \rightarrow R_2 \rightarrow W_2$

Overview of Pre-threaded Concurrent Server



Pre-threaded Concurrent Server

```

sbuf_t sbuf; /* Shared buffer of connected descriptors */

int main(int argc, char **argv) {
    int i, listenfd, connfd;
    socklen_t clientlen;
    struct sockaddr_storage clientaddr;
    pthread_t tid;

    listenfd = Open_listenfd(argv[1]);
    sbuf_init(&sbuf, SBUFSIZE);
    for ( i = 0; i < NTHREADS; i++ ) /* Create a pool of worker threads */
        Pthread_create(&tid, NULL, thread, NULL);
    while (1) {
        clientlen = sizeof(struct sockaddr_storage);
        connfd = Accept(listenfd, (SA *) &clientaddr, &clientlen);
        sbuf_insert(&sbuf, connfd); /* Insert connfd in buffer */
    }
}

void *thread(void *vargp) {
    Pthread_detach(pthread_self());
    while ( 1 ) {
        int connfd = sbuf_remove(&sbuf); /* Remove connfd from buf */
        echo_cnt(connfd);                /* Service client */
        Close(connfd);
    }
}

```

echoservt_pre.c

Pre-threaded Concurrent Server (Cont)

```
static int byte_cnt; /* Byte counter */
static sem_t mutex; /* and the mutex that protects it */
static void init_echo_cnt(void)
{
    Sem_init(&mutex, 0, 1);
    byte_cnt = 0;
}
void echo_cnt(int connfd)
{
    int n;
    char buf[MAXLINE];
    rio_t rio;
    static pthread_once_t once = PTHREAD_ONCE_INIT;

    Pthread_once(&once, init_echo_cnt);
    Rio_readinitb(&rio, connfd);
    while ( ( n = Rio_readlineb(&rio, buf, MAXLINE) ) != 0 ) {
        P(&mutex);
        byte_cnt += n;
        printf("thread %d received %d (%d total) bytes on fd %d\n",
               (int) pthread_self(), n, byte_cnt, connfd);
        V(&mutex);
        Rio_writen(connfd, buf, n);
    }
}
```

echo_cnt.c

Crucial Concept: Thread Safety

- Functions called from a thread must be *thread-safe*
- A function is *thread-safe* iff it will always produce correct results when called repeatedly from multiple concurrent threads
- **Classes of thread-unsafe functions:**
 - Class 1: Functions that do not protect shared variables
 - Class 2: Functions that keep state across multiple invocations
 - Class 3: Functions that return a pointer to a static variable
 - Class 4: Functions that call thread-unsafe functions 😊

Thread-Unsafe Functions (Class 1)

■ Failing to protect shared variables

- Fix: Use P and V semaphore operations
- Example: `goodcnt.c`
- Issue: Synchronization operations will slow down code

Thread-Unsafe Functions (Class 2)

- Relying on persistent state across multiple function invocations
 - Example: Random number generator that relies on static state

```
static unsigned int next = 1;

/* rand: return pseudo-random integer on 0..32767 */
int rand(void)
{
    next = next*1103515245 + 12345;
    return (unsigned int) (next/65536) % 32768;
}

/* srand: set seed for rand() */
void srand(unsigned int seed)
{
    next = seed;
}
```

Thread-Safe Random Number Generator

- Pass state as part of argument
 - and, thereby, eliminate global state

```
/* rand_r - return pseudo-random integer on 0..32767 */  
  
int rand_r(int *nextp)  
{  
    *nextp = *nextp * 1103515245 + 12345;  
    return (unsigned int) (*nextp/65536) % 32768;  
}
```

- Consequence: programmer using `rand_r` must maintain seed

Thread-Unsafe Functions (Class 3)

- Returning a pointer to a static variable
- **Fix 1. Rewrite function so caller passes address of variable to store result**
 - Requires changes in caller and callee
- **Fix 2. Lock-and-copy**
 - Requires simple changes in caller (and none in callee)
 - However, caller must free memory

```
/* lock-and-copy version */  
  
char *ctime_ts(const time_t *timep,  
               char *privatep)  
{  
    char *sharedp;  
  
    P(&mutex);  
    sharedp = ctime(timep);  
    strcpy(privatep, sharedp);  
    V(&mutex);  
  
    return privatep;  
}
```

Thread-Unsafe Functions (Class 4)

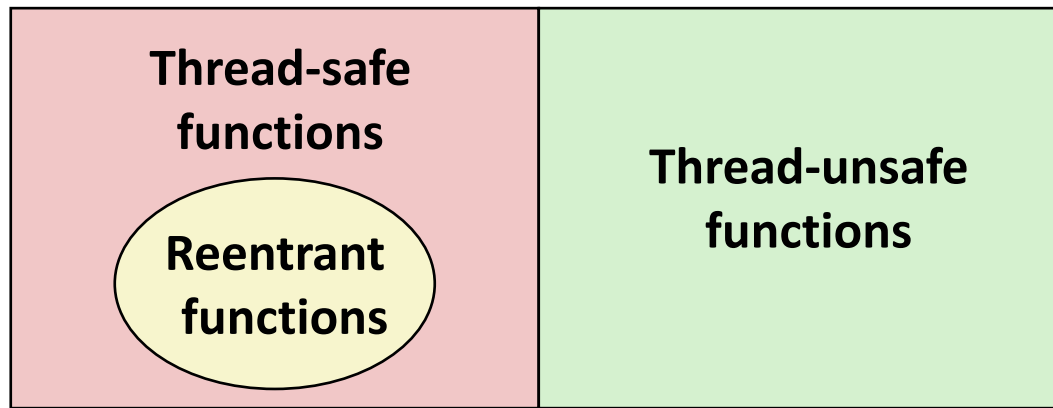
■ Calling thread-unsafe functions

- Calling one thread-unsafe function makes the entire function that calls it thread-unsafe
- Fix: Modify the function so it calls only thread-safe functions 😊

Reentrant Functions

- A function is *reentrant* iff it accesses no shared variables when called by multiple threads
 - Important subset of thread-safe functions
 - Require no synchronization operations
 - Only way to make a Class 2 function thread-safe is to make it *reentrant* (e.g., `rand_r`)

All functions



Examples

■ Not thread-safe, not re-entrant

```
int tmp;  
int add(int a) {  
    tmp = a;  
    return tmp + 10;  
}
```

■ Thread-safe, not re-entrant

```
thread_local int tmp;  
int add(int a) {  
    tmp = a;  
    return tmp + 10;  
}
```

■ Not thread-safe, re-entrant

```
int tmp;  
int add(int a) {  
    tmp = a;  
    return a + 10;  
}
```

■ Thread-safe, re-entrant

```
int add(int a) {  
    return a + 10;  
}
```


Thread-Safe Library Functions

- All functions in the Standard C Library are thread-safe
 - Examples: `malloc`, `free`, `printf`, `scanf`
- Most Unix system calls are thread-safe, with a few exceptions:

Thread-unsafe function	Class	Reentrant version
<code>asctime</code>	3	<code>asctime_r</code>
<code>ctime</code>	3	<code>ctime_r</code>
<code>gethostbyaddr</code>	3	<code>gethostbyaddr_r</code>
<code>gethostbyname</code>	3	<code>gethostbyname_r</code>
<code>inet_ntoa</code>	3	(none)
<code>localtime</code>	3	<code>localtime_r</code>
<code>rand</code>	2	<code>rand_r</code>

One Worry: Races

- A *race* occurs when correctness of the program depends on one thread reaching point x before another thread reaches point y

```
/* A threaded program with a race */
int main()
{
    pthread_t tid[N];
    int i;

    for ( i = 0; i < N; i++ )
        Pthread_create(&tid[i], NULL, thread, &i);
    for ( i = 0; i < N; i++ )
        Pthread_join(tid[i], NULL);
    exit(0);
}

void *thread(void *vargp)
{
    int myid = *((int *)vargp);
    printf("Hello from thread %d\n", myid);
    return NULL;
}
```

N threads are sharing *i*

race.c

Why Race?

```
for ( i = 0; i < N; i++ )  
    Pthread_create(&tid[i], NULL, thread, &i);
```

Main thread

$i = 0$

$i = 1$

Peer thread 0

`myid = *((int *)vargp)`

Race!

- Race between increment of i in main thread and dereference of `vargp` in peer thread
 - If dereference happens while $i = 0$, then OK
 - Otherwise, peer thread gets wrong id value

Could This Race Really Occur?

Main thread

```
int i;

for ( i = 0; i < 100; i++ ) {
    Pthread_create(&tid, NULL,
                  thread, &i);
}
```

Peer thread

```
void *thread(void *vargp) {
    Pthread_detach(pthread_self());
    int i = *((int *)vargp);
    save_value(i);
    return NULL;
}
```


race.c

■ Race test


- If no race, then each thread would get *different* value of *i*
- Set of saved values would consist of one copy each of 0 through 99

Experimental Results


No Race



Single core laptop



Multicore server



The race can really happen!

Race Elimination

```
/* Threaded program without the race */
int main()
{
    pthread_t tid[N];
    int i, *ptr;

    for ( i = 0; i < N; i++ ) {
        ptr = Malloc(sizeof(int));
        *ptr = i;
        Pthread_create(&tid[i], NULL, thread, ptr);
    }
    for ( i = 0; i < N; i++ )
        Pthread_join(tid[i], NULL);
    exit(0);
}


void *thread(void *vargp)
{
    int myid = *((int *)vargp);
    Free(vargp);
    printf("Hello from thread %d\n", myid);
    return NULL;
}
```

Avoid unintended sharing of state

norace.c

Another Worry: Deadlock

- A process is **deadlocked** iff it is waiting for a condition that will never be true
- **Typical scenario**
 - Processes 1 and 2 needs two resources (A and B) to proceed
 - Process 1 acquires A, waits for B
 - Process 2 acquires B, waits for A
 - Both will wait forever!



Deadlock with Semaphores

```
int main()
{
    pthread_t tid[2];
    Sem_init(&mutex[0], 0, 1); /* mutex[0] = 1 */
    Sem_init(&mutex[1], 0, 1); /* mutex[1] = 1 */
    Pthread_create(&tid[0], NULL, count, (void*) 0);
    Pthread_create(&tid[1], NULL, count, (void*) 1);
    Pthread_join(tid[0], NULL);
    Pthread_join(tid[1], NULL);
    printf("cnt=%d\n", cnt);
    exit(0);
}
```

```
void *count(void *vargp)
{
    int i;
    int id = (int) vargp;
    for ( i = 0; i < NITERS; i++ ) {
        P(&mutex[id]); P(&mutex[1-id]);
        cnt++;
        V(&mutex[id]); V(&mutex[1-id]);
    }
    return NULL;
}
```

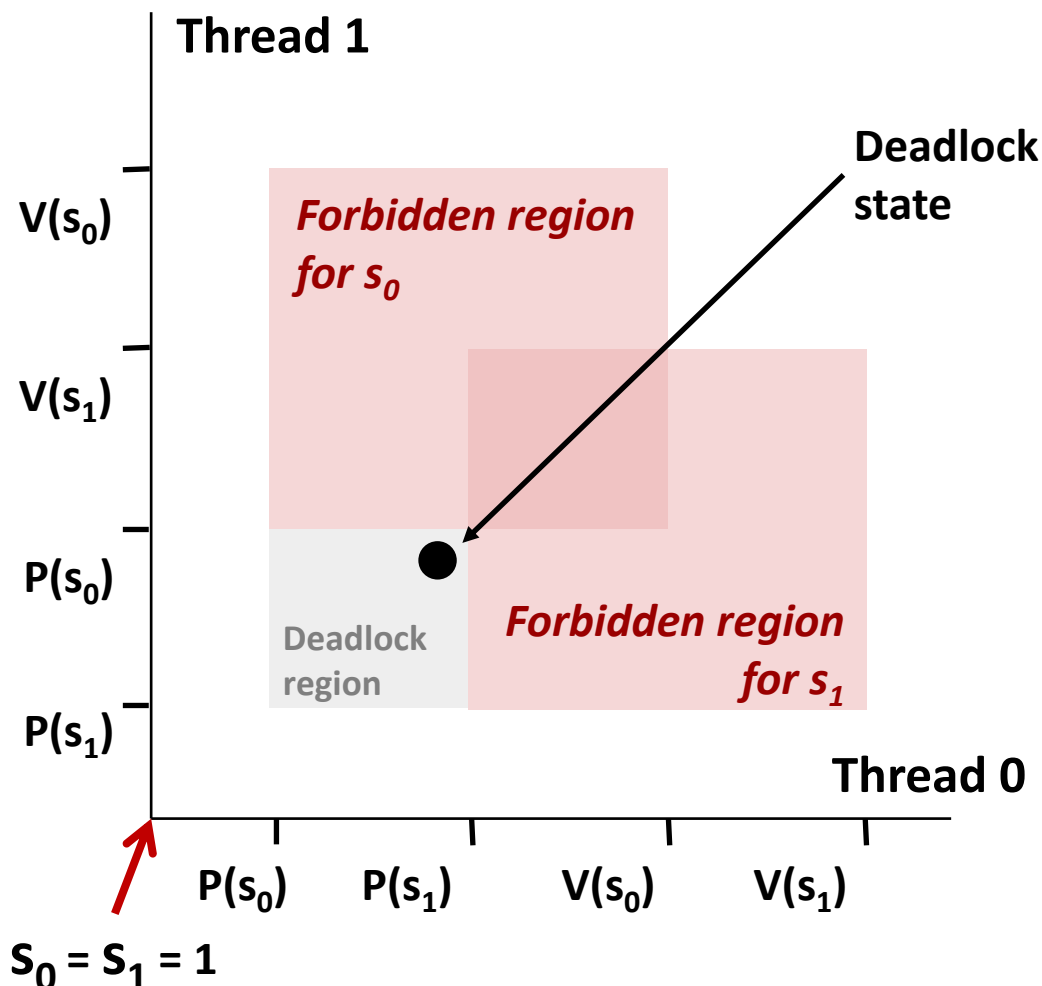
Tid[0]

P(s₀);
P(s₁);
cnt++;
V(s₀);
V(s₁);

Tid[1]

P(s₁);
P(s₀);
cnt++;
V(s₁);
V(s₀);

Deadlock Visualized in Progress Graph



- Locking introduces the potential for **deadlock**: waiting for a condition that will never be true
- Any trajectory that enters the **deadlock region** will eventually reach the **deadlock state**, waiting for either s_0 or s_1 to become nonzero
- Other trajectories luck out and skirt the deadlock region
- Unfortunate fact: deadlock is often nondeterministic (race)

Avoiding Deadlock

```
int main()
{
    pthread_t tid[2];
    Sem_init(&mutex[0], 0, 1); /* mutex[0] = 1 */
    Sem_init(&mutex[1], 0, 1); /* mutex[1] = 1 */
    Pthread_create(&tid[0], NULL, count, (void*) 0);
    Pthread_create(&tid[1], NULL, count, (void*) 1);
    Pthread_join(tid[0], NULL);
    Pthread_join(tid[1], NULL);
    printf("cnt=%d\n", cnt);
    exit(0);
}
```

Acquire shared resources in same order

```
void *count(void *vargp)
{
    int i;
    int id = (int) vargp;
    for ( i = 0; i < NITERS; i++ ) {
        P(&mutex[0]); P(&mutex[1]);
        cnt++;
        V(&mutex[id]); V(&mutex[1-id]);
    }
    return NULL;
}
```

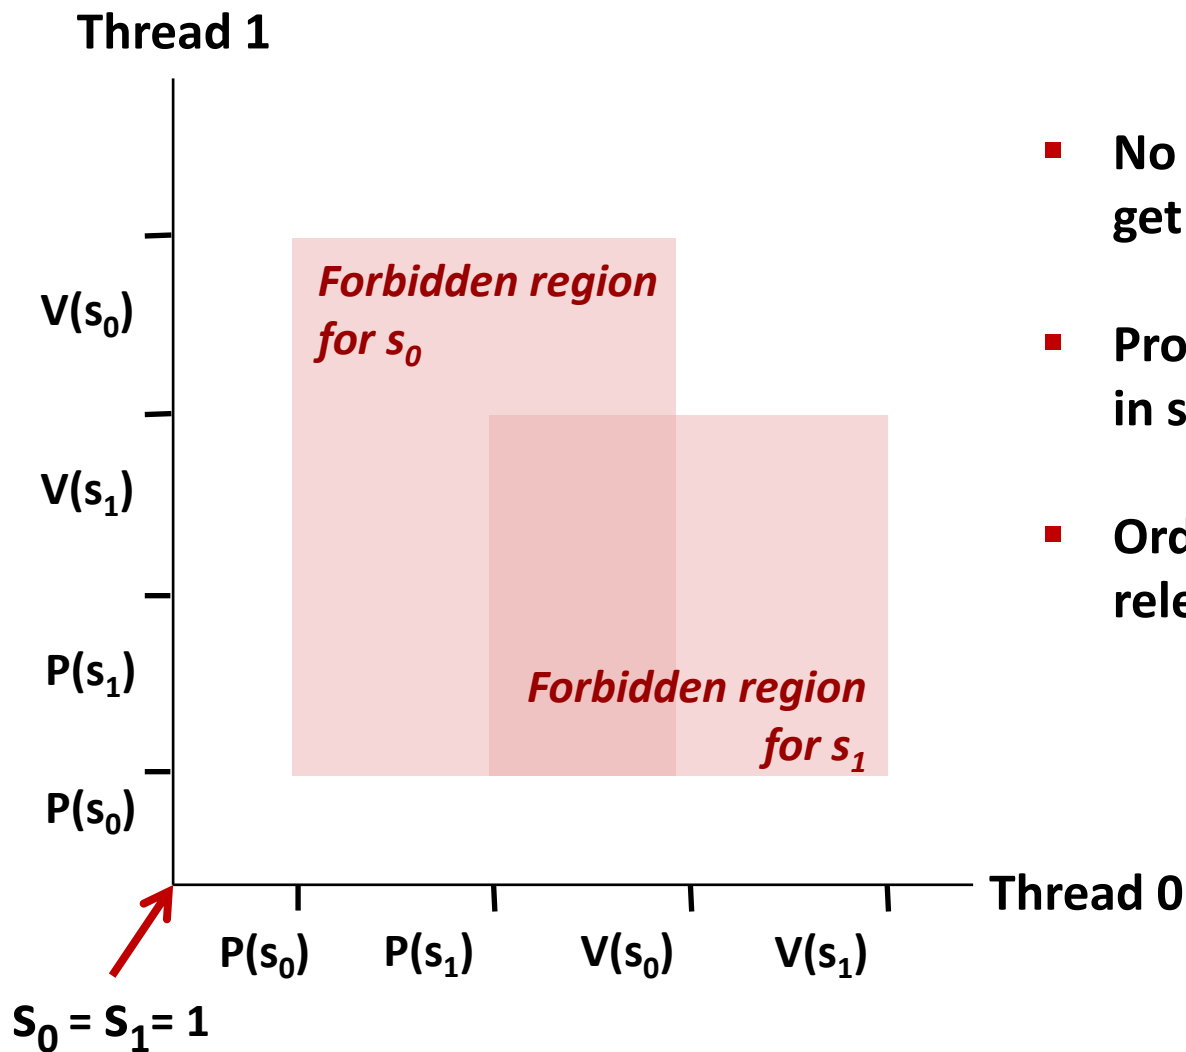
Tid[0]

P(s0);
P(s1);
cnt++;
V(s0);
V(s1);

Tid[1]

P(s0);
P(s1);
cnt++;
V(s1);
V(s0);

Avoided Deadlock in Progress Graph



- No way for trajectory to get stuck
- Processes acquire locks in same order
- Order in which locks released immaterial

Summary

- **Programmers need a clear model of how variables are shared by threads**
- **Variables shared by multiple threads must be protected to ensure mutually exclusive access**
- **Semaphores are a fundamental mechanism for enforcing mutual exclusion**