

Dynamic Memory Allocation

CSE4100: Multicore Programming

Sungyong Park (PhD)

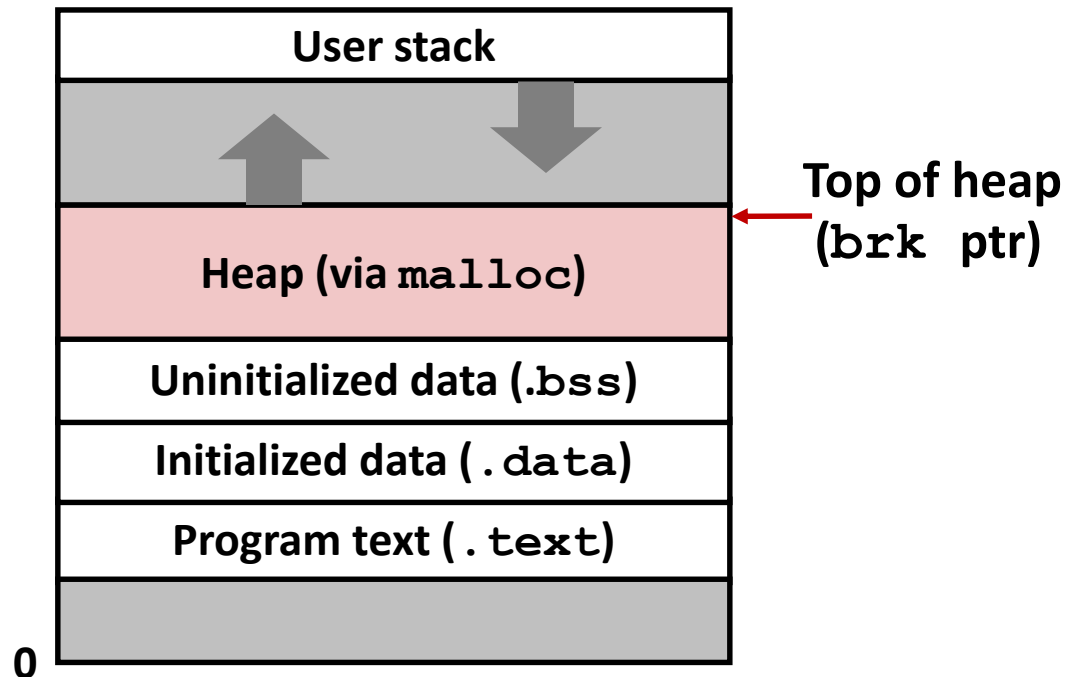

Data-Intensive Computing and Systems Laboratory (DISCOS)

<https://discos.sogang.ac.kr>

Office: R908A, E-mail: parksy@sogang.ac.kr

Dynamic Memory Allocation

- Programmers use *dynamic memory allocators* (such as `malloc`) to acquire virtual memory at run time
 - For data structures whose size is only known at runtime
- Dynamic memory allocators manage an area of process virtual memory known as the *heap*



Dynamic Memory Allocation (Cont)

- Allocator maintains heap as collection of variable sized memory *blocks*, which are either *allocated* or *free*
- Types of allocators
 - *Explicit allocator*: application allocates and frees space
 - e.g., `malloc` and `free` in C
 - *Implicit allocator*: application allocates, but does not free space
 - e.g., garbage collection in Java, ML, and Lisp
- Will focus on simple explicit memory allocation

The malloc Package

```
#include <stdlib.h>
```

```
■ void *malloc(size_t size)
```

- Successful:

- Returns a pointer to a memory block of at least **size** bytes aligned to an **8-byte (x86)** or **16-byte (x86-64)** boundary
- If **size == 0**, returns NULL

- Unsuccessful: returns NULL (0) and sets **errno**

```
■ void free(void *p)
```

- Returns the block pointed at by **p** to pool of available memory
- **p** must come from a previous call to **malloc**, **calloc**, or **realloc**

```
■ Other functions
```

- **calloc**: Version of **malloc** that initializes allocated block to zero
- **realloc**: Changes the size of a previously allocated block
- **sbrk**: Used internally by allocators to grow or shrink the heap

malloc Example

```
#include <stdio.h>
#include <stdlib.h>

void foo(int n) {
    int i, *p;


    /* Allocate a block of n ints */
    p = (int *) malloc(n * sizeof(int));
    if ( p == NULL ) {
        perror("malloc");
        exit(0);
    }

    /* Initialize allocated block */
    for ( i = 0; i < n; i++ )
        p[i] = i;

    /* Return allocated block to the heap */
    free(p);
}
```

Assumptions Made

- **1 square = 1 “word” = 4 bytes**
 - *Word* is the number of bits processed by CPU at the same time
- **Alignment = 8-byte alignment**
 - Alignment depends on whether the code is compiled to run in 32-bit mode (8-byte alignment) or 64-bit mode (16-byte alignment)
- **Example**




Allocation Example

SIZE: word size


`p1 = malloc(4*SIZE)`




`p2 = malloc(5*SIZE)`



`p3 = malloc(6*SIZE)`



`free(p2)`



`p4 = malloc(2*SIZE)`



Constraints

■ Applications

- Can issue arbitrary sequence of **malloc** and **free** requests
- **free** request must be to a **malloc**'d block

■ Explicit allocators

- Can't control the number or size of allocated blocks
- Must respond immediately to **malloc** requests
 - *i.e.*, can't reorder or buffer requests
- Must allocate blocks from free memory
 - *i.e.*, can only place allocated blocks in free memory
- Must align blocks so that they satisfy all alignment requirements
 - 8-byte (x86) or 16-byte (x86-64) alignment
- Can manipulate and modify only free memory
- Can't move the allocated blocks once they are **malloc**'d
 - *i.e.*, compaction is not allowed

Performance Goal: Throughput

- Given some sequence of `malloc` and `free` requests
 - $R_0, R_1, \dots, R_k, \dots, R_{n-1}$
- Goals: maximize *throughput* and *peak memory utilization*
 - These goals are often conflicting
- Throughput
 - Number of completed requests per unit time
 - Example:
 - 5,000 `malloc` calls and 5,000 `free` calls in 10 seconds
 - Throughput is 1,000 operations/second

Performance Goal: Peak Memory Utilization


- Given some sequence of `malloc` and `free` requests
 - $R_0, R_1, \dots, R_k, \dots, R_{n-1}$
- **Def: Aggregate payload P_k**
 - `malloc(p)` results in a block with a **payload** of `p` bytes
 - After request R_k has completed, the **aggregate payload** P_k is the sum of currently allocated payloads
- **Def: Current heap size H_k**
 - Assume H_k is monotonically non-decreasing
 - i.e., heap only grows when allocator uses `sbrk`
- **Def: Peak memory utilization after $k+1$ requests**
 - $U_k = (\max_{i \leq k} P_i) / H_k$

Fragmentation

- Poor memory utilization caused by *fragmentation*
 - *Internal* fragmentation
 - *External* fragmentation

Internal Fragmentation

- For a given block, *internal fragmentation* occurs if payload is smaller than block size



- **Caused by**
 - Overhead of maintaining heap data structures
 - Padding for alignment purposes
 - Explicit policy decisions
(e.g., to return a big block to satisfy a small request)
- **Depends on the pattern of currently allocated requests**
 - Thus, easy to measure

External Fragmentation


SIZE: word size

- Occurs when there is enough aggregate heap memory, but no single free block is large enough


```
p1 = malloc(4*SIZE)
```



```
p2 = malloc(5*SIZE)
```



```
p3 = malloc(6*SIZE)
```



```
free(p2)
```



```
p4 = malloc(7*SIZE)
```

Oops! (what would happen now?)

- Depends on the pattern of future requests
 - Thus, difficult to measure


Implementation Issues

- **How do we know how much memory to free given just a pointer (i.e., `free(p)`)?**
- **How do we keep track of the free blocks?**
 - What do we do with the extra space when allocating a structure that is smaller than the free block it is placed in?
 - How do we pick a block to use for allocation -- many might fit?
 - How do we reinsert freed block?

Knowing How Much to Free


■ Standard method

- Keep the length (**in bytes**) of a block in the word **preceding** the block
 - This word is often called the **header field** or **header**
 - The length field in the header *may* include the size of header and padding
- Requires an extra word for every allocated block




Keeping Track of Free Blocks

- Method 1: *Implicit list* using length—links all blocks



- Method 2: *Explicit list* among the free blocks using pointers



- Method 3: *Segregated free list*

- Different free lists for different size classes


- Method 4: *Blocks sorted by size*

- Can use a balanced tree (e.g., red-black tree) with pointers within each free block, and the length used as a key

Method 1: Implicit List

- For each block we need both **size** and **allocation status**
 - Could store this information in two words: wasteful!
- **Standard trick**
 - If blocks are aligned, some low-order address bits are always 0
 - **8-byte alignment: low-order 3 bits** are always 0
 - Instead of storing always-0 bits, use 1 bit as an allocated/free flag
 - When reading size word, this bit should be *masked out*

*Format of
allocated and
free blocks*




a = 1: Allocated block

a = 0: Free block

Size: block size

**Payload: application data
(allocated blocks only)**

Implicit List: Example



■ 8-byte alignment

- Requires headers are at *non-aligned* positions and payloads are *aligned*
- May require initial unused word and cause some internal fragmentation
- 1 word (0/1) to mark the end of list

■ Advantage

- Simplicity

■ Disadvantage

- Placing an allocated block requires a search of the list to find a free block (sometimes the whole list)

Implicit List: Finding a Free Block

■ *First fit*

- Search the list from the beginning, choose the *first* free block that fits
- Can take linear time in total number of blocks (allocated and free)
- In practice, it can cause *splinters* at the beginning of list

■ *Next fit*

- Search the list from the position where previous search finished
- Should often be faster than *first fit*: avoids re-scanning unhelpful blocks
- Some research indicates that fragmentation is worse


■ *Best fit*

- Search the list, choose the *best* free block: fits, with the fewest bytes left over
- Keeps fragments small—usually improves memory utilization
- Will typically run slower than first fit
- No guarantee of optimality

Implicit List: Allocating in Free Block

■ Allocating in a free block: *splitting*


- Since allocated space might be smaller than free space, we might want to split the block



Implicit List: Freeing a Block

■ Simplest implementation

- Need only clear the “allocated” flag
- But can lead to “false fragmentation”




SIZE: word size

Oops!

***There is enough contiguous free space,
but the allocator won't be able to find it***

Implicit List: Coalescing

- Join (*coalesce*) with next/previous blocks, if they are free
 - Coalescing with next block




- How do we coalesce with *previous* block?
 - How do we know where it starts?
 - How can we determine whether its allocated?


Implicit List: Bidirectional Coalescing

■ *Boundary tags* [Knuth73]

- Replicate size/allocated word at *bottom (end)* of free blocks
- Allows us to traverse the list *backwards*, but requires extra space
- Important and general technique!



*Format of
allocated and
free blocks*

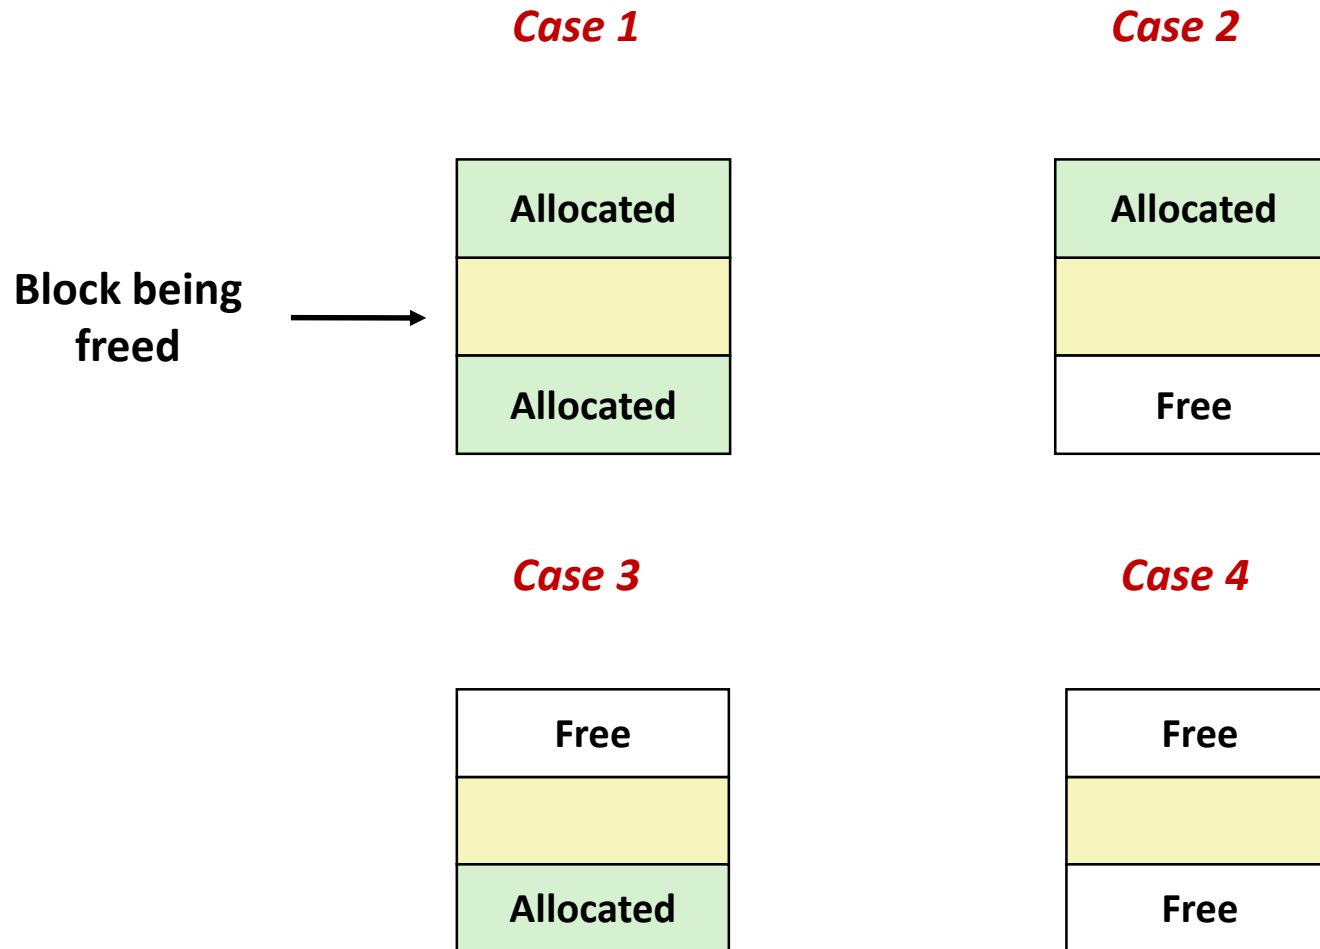


a = 1: Allocated block
a = 0: Free block

Size: Total block size


Payload: Application data
(allocated blocks only)

Constant Time (Immediate) Coalescing




Constant Time (Immediate) Coalescing (Cont)


Case 1




Case 2



Case 3



Case 4



Disadvantages of Boundary Tags

- **Extra space overhead**
- **Which blocks need the footer tag?**
 - If the previous block is not free (allocated), we don't need the boundary tag of the previous block
 - Otherwise, we need the boundary tag of the previous block, because it should be able to tell the size of the block
- **Can it be optimized?**
 - Store the allocated/free bit of the previous block in one of the excess low-order bits of the current block
 - Then, allocated blocks would not need footers and the extra space can be used for payload

Summary of Key Allocator Policies

■ Placement policy

- First-fit, next-fit, best-fit, etc.
- Trades off lower throughput for less fragmentation
- *Interesting observation*: segregated free lists (covered later) approximate a best fit placement policy without having to search entire free list

■ Splitting policy

- When do we go ahead and split free blocks?
- How much internal fragmentation are we willing to tolerate?

■ Coalescing policy

- *Immediate coalescing*: coalesce each time **free** is called
- *Deferred coalescing*: try to improve performance of **free** by deferring coalescing until needed
 - Coalesce as you scan the free list for **malloc**
 - Coalesce when the amount of external fragmentation reaches some threshold


Implementing a Simple Allocator

■ Export three functions to application programs

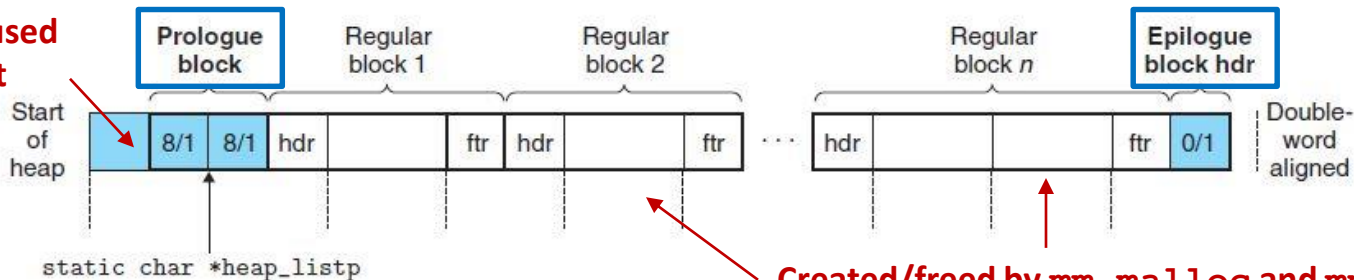
- `int mm_init(void)`
- `void *mm_malloc(size_t size) -> same as malloc()`
- `void *mm_free(void *ptr) -> same as free()`

■ Assumption

- Use *implicit list* with *boundary tag* (right)
- Use **4-byte** word and **8-byte** alignment
- Minimum block size is **16 bytes** (4 words) due to alignment and block format
- Heap consists of 1 prologue block, 0 or more regular blocks and 1 epilogue block



Intentionally unused
for alignment



Implementing a Simple Allocator (Cont)

■ Useful constants and macros

code/vm/malloc/mm.c

```

1  /* Basic constants and macros */
2  #define WSIZE      4      /* Word and header/footer size (bytes) */
3  #define DSIZE     8      /* Double word size (bytes) */
4  #define CHUNKSIZE (1<<12) /* Extend heap by this amount (bytes) */
5
6  #define MAX(x, y) ((x) > (y)? (x) : (y))
7
8  /* Pack a size and allocated bit into a word */
9  #define PACK(size, alloc) ((size) | (alloc))
10
11 /* Read and write a word at address p */
12 #define GET(p)      (*(unsigned int *) (p))
13 #define PUT(p, val) (*(unsigned int *) (p) = (val))
14
15 /* Read the size and allocated fields from address p */
16 #define GET_SIZE(p) (GET(p) & ~0x7) ← 8-byte alignment
17 #define GET_ALLOC(p) (GET(p) & 0x1) ← LSB is alloc/free bit
18
19 /* Given block ptr bp, compute address of its header and footer */
20 #define HDRP(bp)    ((char *) (bp) - WSIZE)
21 #define FTRP(bp)    ((char *) (bp) + GET_SIZE(HDRP(bp)) - DSIZE)
22
23 /* Given block ptr bp, compute address of next and previous blocks */
24 #define NEXT_BLKP(bp) ((char *) (bp) + GET_SIZE(((char *) (bp) - WSIZE)))
25 #define PREV_BLKP(bp) ((char *) (bp) - GET_SIZE(((char *) (bp) - DSIZE)))

```

code/vm/malloc/mm.c

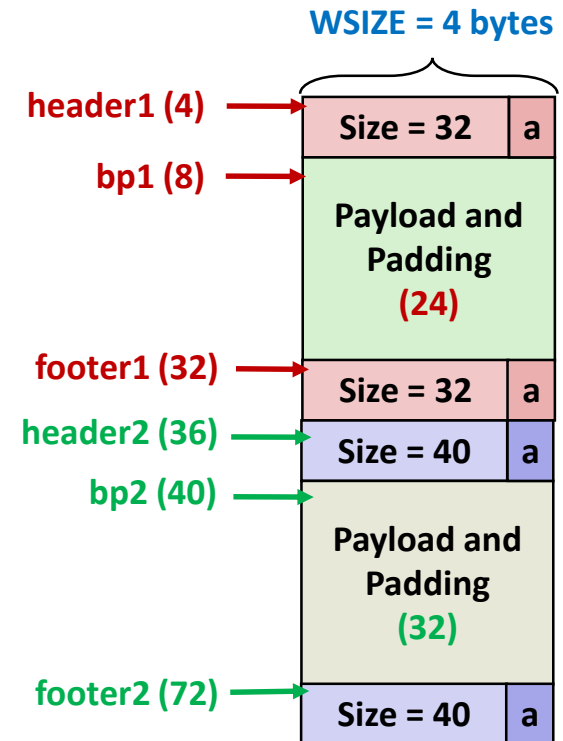


Figure 9.43 Basic constants and macros for manipulating the free list.

Implementing a Simple Allocator (Cont)

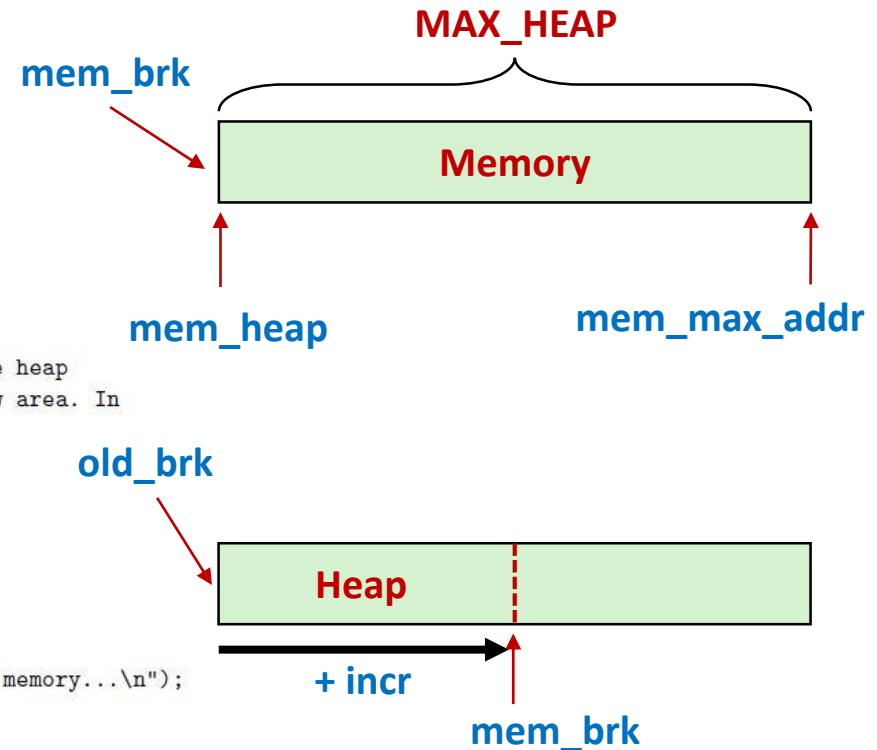
■ Memory initialization for heap and heap extension

code/vm/malloc/memlib.c

```

1  /* Private global variables */
2  static char *mem_heap;    /* Points to first byte of heap */
3  static char *mem_brk;     /* Points to last byte of heap plus 1 */
4  static char *mem_max_addr; /* Max legal heap addr plus 1 */
5
6  /*
7   * mem_init - Initialize the memory system model
8   */
9  void mem_init(void)
10 {
11     mem_heap = (char *)Malloc(MAX_HEAP);
12     mem_brk = (char *)mem_heap;
13     mem_max_addr = (char *) (mem_heap + MAX_HEAP);
14 }
15
16 /*
17 * mem_sbrk - Simple model of the sbrk function. Extends the heap
18 *   by incr bytes and returns the start address of the new area. In
19 *   this model, the heap cannot be shrunk.
20 */
21 void *mem_sbrk(int incr)
22 {
23     char *old_brk = mem_brk;
24
25     if ( (incr < 0) || ((mem_brk + incr) > mem_max_addr) ) {
26         errno = ENOMEM;
27         fprintf(stderr, "ERROR: mem_sbrk failed. Ran out of memory...\n");
28         return (void *)-1;
29     }
30     mem_brk += incr;
31     return (void *)old_brk;
32 }

```

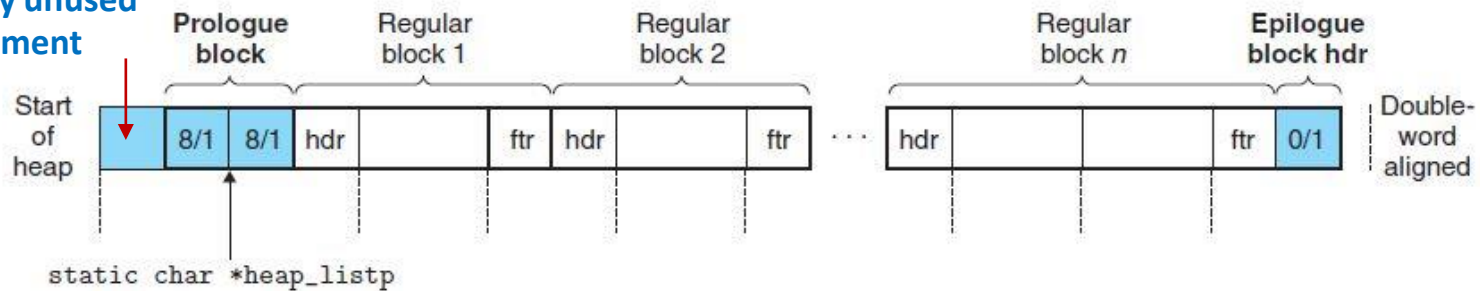


code/vm/malloc/memlib.c

Implementing a Simple Allocator (Cont)

■ Heap initialization

Intentionally unused
for alignment



code/vm/malloc/mm.c

```

1  int mm_init(void)
2  {
3      /* Create the initial empty heap */
4      if ((heap_listp = mem_sbrk(4*WSIZE)) == (void *)-1)
5          return -1;
6      PUT(heap_listp, 0);                          /* Alignment padding */
7      PUT(heap_listp + (1*WSIZE), PACK(DSIZE, 1)); /* Prologue header */
8      PUT(heap_listp + (2*WSIZE), PACK(DSIZE, 1)); /* Prologue footer */
9      PUT(heap_listp + (3*WSIZE), PACK(0, 1));      /* Epilogue header */
10     heap_listp += (2*WSIZE);
11
12     /* Extend the empty heap with a free block of CHUNKSIZE bytes */
13     if (extend_heap(CHUNKSIZE/WSIZE) == NULL)
14         return -1;
15     return 0;
16 }
  
```

Create an empty heap with only prologue and epilogue blocks (16 bytes – shaded in blue)

Initially, create a 1024 word heap (4096 bytes)

code/vm/malloc/mm.c

Figure 9.44 mm_init creates a heap with an initial free block.

Implementing a Simple Allocator (Cont)

■ Extend a heap

```

1  static void *extend_heap(size_t words)
2  {
3      char *bp;
4      size_t size;
5
6      /* Allocate an even number of words to maintain alignment */
7      size = (words % 2) ? (words+1) * WSIZE : words * WSIZE;
8      if ((long)(bp = mem_sbrk(size)) == -1)
9          return NULL;
10
11     /* Initialize free block header/footer and the epilogue header */
12     PUT(HDRP(bp), PACK(size, 0));          /* Free block header */
13     PUT(FTRP(bp), PACK(size, 0));          /* Free block footer */
14     PUT(HDRP(NEXT_BLKP(bp)), PACK(0, 1)); /* New epilogue header */
15
16     /* Coalesce if the previous block was free */
17     return coalesce(bp);
18 }

```

code/vm/malloc/mm.c

code/vm/malloc/mm.c


Figure 9.45 `extend_heap` extends the heap with a new free block.

Implementing a Simple Allocator (Cont)


■ Heap allocation

If requested size ≤ 8 bytes,
allocate the minimum block size
(16 bytes) to include
header and footer

4 bytes



mm_malloc(7)



mm_malloc(9)

code/vm/malloc/mm.c

```

1 void *mm_malloc(size_t size)
2 {
3     size_t asize; /* Adjusted block size */
4     size_t extendsize; /* Amount to extend heap if no fit */
5     char *bp;
6
7     /* Ignore spurious requests */
8     if (size == 0)
9         return NULL;
10
11     /* Adjust block size to include overhead and alignment reqs. */
12     if (size <= DSIZE)
13         asize = 2*DSIZE;
14     else
15         asize = DSIZE * ((size + (DSIZE) + (DSIZE-1)) / DSIZE);
16
17     /* Search the free list for a fit */
18     if ((bp = find_fit(asize)) != NULL) {
19         place(bp, asize);
20         return bp;
21     }
22
23     /* No fit found. Get more memory and place the block */
24     extendsize = MAX(asize, CHUNKSIZE);
25     if ((bp = extend_heap(extendsize/WSIZE)) == NULL)
26         return NULL;
27     place(bp, asize);
28     return bp;
29 }

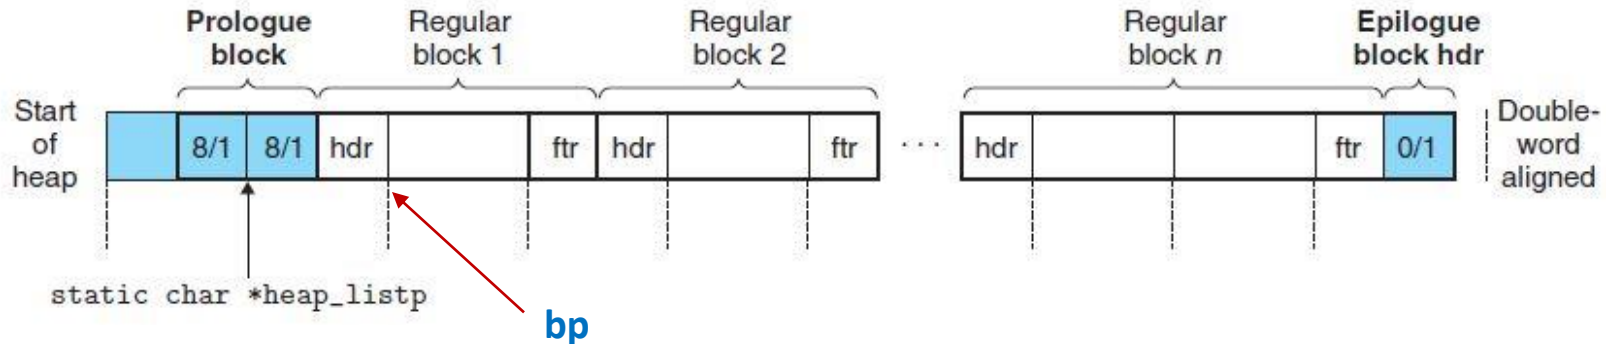
```

code/vm/malloc/mm.c

Figure 9.47 mm_malloc allocates a block from the free list.

Implementing a Simple Allocator (Cont)

■ Search the free list (*first-fit*)



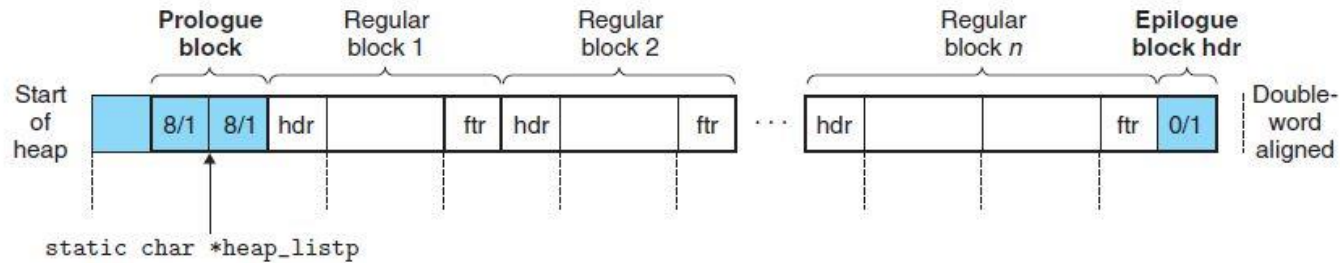
```

code/vm/malloc/mm.c
1  static void *find_fit(size_t asize)
2  {
3      /* First-fit search */
4      void *bp;
5
6      for (bp = heap_listp; GET_SIZE(HDRP(bp)) > 0; bp = NEXT_BLKBP(bp)) {
7          if (!GET_ALLOC(HDRP(bp)) && (asize <= GET_SIZE(HDRP(bp)))) {
8              return bp;
9          }
10     }
11     return NULL; /* No fit */
12 #endif
13 }
code/vm/malloc/mm.c

```

Implementing a Simple Allocator (Cont)

■ Split the free list if needed



```

1  static void place(void *bp, size_t asize)
2  {
3      size_t csize = GET_SIZE(HDRP(bp));
4
5      if ((csize - asize) >= (2*DSIZE)) {
6          PUT(HDRP(bp), PACK(asize, 1));
7          PUT(FTRP(bp), PACK(asize, 1));
8          bp = NEXT_BLK(bp);
9          PUT(HDRP(bp), PACK(csize-asize, 0));
10         PUT(FTRP(bp), PACK(csize-asize, 0));
11     }
12     else {
13         PUT(HDRP(bp), PACK(csize, 1));
14         PUT(FTRP(bp), PACK(csize, 1));
15     }
16 }

```

code/vm/malloc/mm.c

If enough space is left,
split the list and update metadata

If left-over size is less than
the minimum block size (16 bytes),
ignore (internal fragmentation)

code/vm/malloc/mm.c


Implementing a Simple Allocator (Cont)

Free allocated heap

```

1 void mm_free(void *bp)
2 {
3     size_t size = GET_SIZE(HDRP(bp));
4     PUT(HDRP(bp), PACK(size, 0));
5     PUT(FTRP(bp), PACK(size, 0));
6     coalesce(bp);
7 }
8
9

```



```

10 static void *coalesce(void *bp)
11 {
12     size_t prev_alloc = GET_ALLOC(FTRP(PREV_BLKPTR(bp)));
13     size_t next_alloc = GET_ALLOC(HDRP(NEXT_BLKPTR(bp)));
14     size_t size = GET_SIZE(HDRP(bp));
15
16     if (prev_alloc && next_alloc) { /* Case 1 */
17         return bp;
18     }
19
20     else if (prev_alloc && !next_alloc) { /* Case 2 */
21         size += GET_SIZE(HDRP(NEXT_BLKPTR(bp)));
22         PUT(HDRP(bp), PACK(size, 0));
23         PUT(FTRP(bp), PACK(size, 0));
24     }
25
26     else if (!prev_alloc && next_alloc) { /* Case 3 */
27         size += GET_SIZE(HDRP(PREV_BLKPTR(bp)));
28         PUT(FTRP(bp), PACK(size, 0));
29         PUT(HDRP(PREV_BLKPTR(bp)), PACK(size, 0));
30         bp = PREV_BLKPTR(bp);
31     }
32
33     else { /* Case 4 */
34         size += GET_SIZE(HDRP(PREV_BLKPTR(bp))) +
35             GET_SIZE(FTRP(NEXT_BLKPTR(bp)));
36         PUT(HDRP(PREV_BLKPTR(bp)), PACK(size, 0));
37         PUT(FTRP(NEXT_BLKPTR(bp)), PACK(size, 0));
38         bp = PREV_BLKPTR(bp);
39     }
40     return bp;
41 }

```


code/vm/malloc/mm.c

Implicit Lists Summary


- **Implementation: very simple**
- **Allocate cost**
 - linear time worst case
- **Free cost**
 - constant time worst case
 - even with coalescing
- **Memory usage**
 - will depend on placement policy
 - First-fit, next-fit or best-fit
- **Not used in practice for `malloc/free` because of linear-time allocation**
 - used in many special purpose applications
- **However, the concepts of splitting and boundary tag coalescing are general to *all* allocators**

Keeping Track of Free Blocks

- Method 1: *Implicit list* using length—links all blocks




- Method 2: *Explicit list* among the free blocks using pointers




- Method 3: *Segregated free list*
 - Different free lists for different size classes
- Method 4: *Blocks sorted by size*
 - Can use a balanced tree (e.g., Red-Black tree) with pointers within each free block, and the length used as a key

Method 2: Explicit Free Lists

Allocated (as before)



Free




- Maintain list(s) of *free* blocks, not *all* blocks
 - The *next* free block could be anywhere
 - So, we need to store forward/back pointers, not just sizes
 - Still need boundary tags for coalescing
 - Luckily, we track only free blocks, so we can use *payload* area

Explicit Free Lists

■ Logically:




■ Physically: blocks can be in any order



Allocating From Explicit Free Lists


conceptual graphic

Before



After

(with splitting)




Freeing With Explicit Free Lists

- **Insertion policy:** Where in the free list do you put a newly freed block?
- **Unordered**
 - LIFO (last-in-first-out) policy
 - Insert freed block at the **beginning** of the free list
 - FIFO (first-in-first-out) policy
 - Insert freed block at the **end** of the free list
 - **Pro:** simple and constant time
 - **Con:** studies suggest that fragmentation is *worse* than address ordered
- **Address-ordered policy**
 - Insert freed blocks so that free list blocks are always in *address order*
 $addr(prev) < addr(curr) < addr(next)$
 - **Con:** requires search
 - **Pro:** studies suggest that fragmentation is *lower* than LIFO/FIFO

Freeing With a LIFO Policy (Case 1)


conceptual graphic

Before




- Insert the freed block at the root of the list

After




Freeing With a LIFO Policy (Case 2)

conceptual graphic



- Splice out successor block, coalesce both memory blocks and insert the new block at the root of the list




[illegible]


- After*
-
- The diagram illustrates the state of a B-tree after inserting a new leaf. The root node points to a new leaf (green circle) and an existing leaf (red circle). The new leaf is linked to the existing leaf via a green arrow. The existing leaf is linked to the next leaf via a red arrow. The new leaf is linked to the root via a green arrow.

Freeing With a LIFO Policy (Case 4)

conceptual graphic



- Splice out predecessor and successor blocks, coalesce all 3 memory blocks and insert the new block at the root of the list



Explicit List Summary

■ Comparison to implicit list


- Allocate is linear time in number of *free* blocks instead of *all* blocks
 - *Much faster* when most of the memory is full
- Slightly more complicated allocate and free since splicing blocks in and out of the list is needed
- Some extra space for the links (2 extra words needed for each block)
 - Does this increase internal fragmentation?

■ Most common use of linked lists is in conjunction with *segregated free lists*


- Keep multiple linked lists of different size classes, or possibly for different types of objects

Keeping Track of Free Blocks

- Method 1: *Implicit list* using length—links all blocks



- Method 2: *Explicit list* among the free blocks using pointers



- Method 3: *Segregated free list*


- Different free lists for different size classes

- Method 4: *Blocks sorted by size*

- Can use a balanced tree (e.g., Red-Black tree) with pointers within each free block, and the length used as a key

Method 3: Segregated List (Seglist) Allocators

- Each *size class* of blocks has its own free list



- Often have separate classes for each small size
- For larger sizes: One class for each power-of-two (2^n) size

Seglist Allocator

- Given an array of free lists, each one for some size class
- To allocate a block of size n
 - Search appropriate free list for block of size $m > n$
 - If an appropriate block is found
 - Split block and place fragment on appropriate list (optional)
 - If no block is found, try next larger class
 - Repeat until block is found
- If no block is found
 - Request additional heap memory from OS (using **sbrk()**)
 - Allocate block of n bytes from this new memory
 - Place remainder as a single free block in the largest size class

Seglist Allocator (Cont)

■ To free a block

- Coalesce and place on appropriate list

■ Advantages of seglist allocators

- Higher throughput
 - Log time for power-of-two size classes
- Better memory utilization
 - First-fit search of segregated free list approximates a best-fit search of entire heap
 - Extreme case: Giving each block its own size class is equivalent to best-fit

More Info on Allocators

- **D. Knuth, “*The Art of Computer Programming*”, 2nd edition, Addison Wesley, 1973**
 - The classic reference on dynamic storage allocation

- **Wilson et al, “*Dynamic Storage Allocation: A Survey and Critical Review*”, Proc. 1995 Int’l Workshop on Memory Management, Kinross, Scotland, Sept., 1995**
 - Comprehensive survey

Implicit Memory Management

- ***Garbage collection***: automatic reclamation of heap-allocated storage—application never has to free

```
void foo() {  
    int *p = malloc(128);  
    return; /* p block is now garbage */  
}
```

- **Common in many dynamic languages**
 - Python, Ruby, Java, Perl, ML, Lisp, Mathematica
- **Variants (“conservative” garbage collectors) exist for C and C++**
 - However, cannot necessarily collect all garbage

Garbage Collection

- **How does the memory manager know when memory can be freed?**
 - In general, we cannot know when memory is going to be used in the future
 - But, we can tell that certain blocks cannot be used if there are *no pointers* to them

- **Must make certain assumptions about pointers**
 - Memory manager can distinguish pointers from non-pointers
 - All pointers point to the start of a block
 - Cannot hide pointers
(e.g., by type casting them to an `int`, and then back again)


Classical GC Algorithms

- **Mark-and-sweep collection (McCarthy, 1960)**
 - Does not move blocks (unless you also *compact*)
- **Reference counting (Collins, 1960)**
 - Does not move blocks (not discussed)
- **Copying collection (Minsky, 1963)**
 - Moves blocks (not discussed)
- **Generational Collectors (Lieberman and Hewitt, 1983)**
 - Collection based on lifetimes (not discussed)
- **For more information:**
Jones and Lin, *“Garbage Collection: Algorithms for Automatic Dynamic Memory”*, John Wiley & Sons, 1996

Memory as a Graph

■ View memory as a directed graph

- Each block is a *node* in the graph
- Each pointer is an *edge* in the graph
- Locations not in the heap that contain pointers into the heap are called **root nodes** (e.g., registers, variables on the stack, global variables)





A node (block) is **reachable** if there is a path from any root node to that node

Non-reachable nodes are **garbage** (can never be used again by the application)

Mark and Sweep Collection

- Can build on top of `malloc/free` package
 - Allocate using `malloc` until you run out of space
- When we run out of space
 - Use extra **mark bit** in the header of each block
 - **Mark phase**: Start at roots and set mark bit on each reachable block
 - **Sweep phase**: Scan all blocks and free blocks that are not marked



Pseudocode for Mark and Sweep

■ Pseudocode

Call mark function
for each root node
(mark phase)

(a) mark function

```
void mark(ptr p) {
    if ((b = isPtr(p)) == NULL)
        return;
    if (blockMarked(b))
        return;
    markBlock(b);
    len = length(b);
    for (i=0; i < len; i++)
        mark(b[i]);
    return;
}
```

Called once
(sweep phase)

(b) sweep function

```
void sweep(ptr b, ptr end) {
    while (b < end) {
        if (blockMarked(b))
            unmarkBlock(b);
        else if (blockAllocated(b))
            free(b);
        b = nextBlock(b);
    }
    return;
}
```

■ Helper functions

Our description of Mark&Sweep will assume the following functions, where ptr is defined as typedef void *ptr:

ptr isPtr(ptr p). If p points to some word in an allocated block, it returns a pointer b to the beginning of that block. Returns NULL otherwise.

int blockMarked(ptr b). Returns true if block b is already marked.

int blockAllocated(ptr b). Returns true if block b is allocated.

void markBlock(ptr b). Marks block b.

int length(ptr b). Returns the length in words (excluding the header) of block b.


void unmarkBlock(ptr b). Changes the status of block b from marked to unmarked.

ptr nextBlock(ptr b). Returns the successor of block b in the heap.

Conservative Mark & Sweep in C

■ A conservative garbage collector for C programs


- `is_ptr(p)` determines whether `p` is a pointer by checking if it points to an allocated block of memory but it is hard to determine in C
- C pointers can point to the middle of a block



■ How to find the beginning of a block?

- Can use a balanced binary tree to keep track of all allocated blocks (key is start-of-block)
- Balanced-tree pointers can be stored in the header of each allocated block -> `is_ptr(p)` uses this tree to see if `p` falls within the block

Allocated block header



Left: smaller addresses
Right: larger addresses

Memory-Related Perils and Pitfalls

- Dereferencing bad pointers
- Reading uninitialized memory
- Overwriting memory
- Referencing non-existent variables
- Freeing blocks multiple times
- Referencing freed blocks
- Failing to free blocks

C Operators and Precedence

Operators	Associativity
<code>() [] -> . ++ --</code>	left to right
<code>! ~ ++ -- + - * & (type) sizeof</code>	right to left
<code>* / %</code>	left to right
<code>+ -</code>	left to right
<code><< >></code>	left to right
<code>< <= > >=</code>	left to right
<code>== !=</code>	left to right
<code>&</code>	left to right
<code>^</code>	left to right
<code> </code>	left to right
<code>&&</code>	left to right
<code> </code>	left to right
<code>? :</code>	right to left
<code>= += -= *= /= %= &= ^= != <<= >>=</code>	right to left
<code>,</code>	left to right

Source: K&R page 53, updated

- `->`, `()`, and `[]` have high precedence, with `*` and `&` just below
- Unary `+`, `-`, and `*` have higher precedence than binary forms

C Pointer Declarations: Test Yourself!

```
int *p
```

p is a pointer to int

```
int **p
```

p is a pointer to a pointer to an int

```
int *p[4]
```

p is an array[4] of pointer to int

```
int *(p[4])
```

p is an array[4] of pointer to int

```
int (*p)[4]
```

p is a pointer to an array[4] of int

```
int *func()
```

func is a function returning a pointer to int

```
int (*func)()
```

func is a pointer to a function returning int


Source: K&R Sec 5.12

Example


`int *p`




`int **p`




`int *p[4]`



`int *(p[4])`



`int (*p)[4]`



`int *f()`

`int (*f)()`

`int *p = f()`

f is a pointer to a function returning int

Dereferencing Bad Pointers

■ The classic scanf bug

```
int val;  
  
...  
  
scanf ("%d", val) ;
```


Reading Uninitialized Memory

- Assuming that heap data is initialized to zero


```
/* return y = Ax */
int *matvec(int **A, int *x) {
    int *y = malloc(N*sizeof(int));
    int i, j;

    for ( i = 0; i < N; i++ )
        for ( j = 0; j < N; j++ )
            y[i] += A[i][j] * x[j];
    return y;
}
```

Overwriting Memory (Case 1)

- Allocating the (possibly) wrong sized object

```
int **p;  
  
p = malloc(N*sizeof(int)); /* sizeof(int *) */  
  
for ( i = 0; i < N; i++ ) {  
    p[i] = malloc(M*sizeof(int));  
}
```



Return values for `sizeof(int)` and `sizeof(int *)` are different in 64 bit machines!

Overwriting Memory (Case 2)

■ Off-by-one errors

```
char **p;  
  
p = malloc(N*sizeof(int *));  
  
for ( i = 0; i <= N; i++ ) {  
    p[i] = malloc(M*sizeof(int));  
}
```


```
char *p;  
  
p = malloc(strlen(s));  
strcpy(p,s);
```

← Copy including null character

Overwriting Memory (Case 3)

- Not checking the max string size

```
char s[8];  
int i;  
  
gets(s);  /* reads "123456789" from stdin */
```



Basis for classic buffer overflow attacks!

Overwriting Memory (Case 4)

■ Misunderstanding pointer arithmetic

```
int *search(int *p, int val) {  
    while ( *p && *p != val )  
        p += sizeof(int); /* should be p++ */  
    return p;  
}
```

Overwriting Memory (Case 5)

- Referencing a pointer instead of the object it points to

```
int *BinheapDelete(int **binheap, int *size) {
    int *packet;
    packet = binheap[0];
    binheap[0] = binheap[*size - 1];
    *size--; /* This should be (*size)-- */
    Heapify(binheap, *size, 0);
    return(packet);
}
```

Operators

() [] -> . ++ --
 ! ~ ++ -- + - * & (type) sizeof
 * / %
 + -
 << >>
 < <= > >=
 == !=
 &
 ^
 |
 &&
 ||
 ?:
 = += -= *= /= %= &= ^= != <<= >>=
 ,

Associativity

left to right
 right to left
 left to right
 left to right
 left to right
 left to right
 left to right
 left to right
 left to right
 right to left
 right to left
 left to right

Referencing Nonexistent Variables

- Forgetting that local variables disappear when a function returns

```
int *foo () {  
    int val;  
  
    return &val;  
}
```

Freeing Blocks Multiple Times

■ Nasty!

```
x = malloc(N*sizeof(int));  
    <manipulate x>  
free(x);  
  
y = malloc(M*sizeof(int));  
    <manipulate y>  
free(x);
```


Referencing Freed Blocks

■ Evil!

```
x = malloc(N*sizeof(int));  
  <manipulate x>  
free(x);  
  ...  
y = malloc(M*sizeof(int));  
for ( i = 0; i < M; i++ )  
    y[i] = x[i]++;
```

Failing to Free Blocks (Memory Leaks)

- Slow, long-term killer!

```
foo() {  
    int *x = malloc(N*sizeof(int));  
    ...  
    return; /* x is garbage at this point */  
}
```

Dealing With Memory Bugs

■ Debugger: **gdb** (<https://www.sourceware.org/gdb/>)

- Good for finding bad pointer dereferences
- Hard to detect the other memory bugs

■ Binary translator: **valgrind** (<https://valgrind.org/>)

- Powerful debugging and analysis technique
- Rewrites text section of executable object file
- Checks each individual reference at runtime
 - Bad pointers, overwrites, references outside of allocated block

■ **glibc malloc** contains checking code

- `setenv MALLOC_CHECK_ 3`
- `int mallopt(int param, int value)`